

Statistical Learning Theory Notes

Nong Minh Hieu¹

¹ School of Physical and Mathematical Sciences, Nanyang Technological University (NTU - Singapore)

Contents

1	Probability settings	2
1.1	Classification problem	2
1.2	Goal of classification	4
2	Bayes classifier	5
2.1	Properties of Bayes Risk	5
2.2	Likelihood Ratio Test	7
2.3	Plug-in classifier	8
2.4	End of chapter exercises	10
3	Hoeffding's inequality	13
3.1	Markov's Inequality	13
3.2	Hoeffding's Inequality	14
3.3	Convergence of Empirical Risk	15
3.4	KL-divergence & Hypothesis Testing	16
3.5	End of chapter exercises	19
4	Empirical Risk Minimization	21
4.1	Uniform Deviation Bounds	21
4.2	PAC Learning & Sample Complexity	24
4.3	Zero-error case	24
4.4	End of chapter exercises	27
A	Related topics	28
A.1	Neyman-Pearson Lemma	28
A.1.1	Type I & Type II errors	28
A.1.2	Neyman-Pearson Lemma	29
B	List of Definitions	30
C	Important Theorems	30
D	Important Corollaries	30
E	Important Propositions	30
F	References	31

1 Probability settings

1.1 Classification problem

Definition 1.1 (Classifier (h)).

In **classification problems**, we consider pairs (x, y) where $x \in \mathcal{X}$ and $y \in \mathcal{Y}$. Where:

- \mathcal{X} is the space of **feature vectors**.
- \mathcal{Y} is the space of **labels**.

A classifier is a function $h : \mathcal{X} \rightarrow \mathcal{Y}$ which aims to assign correct labels to given feature vectors.

Remark : The key assumptions of classification problems are:

- There exists a joint distribution P_{XY} on $\mathcal{X} \times \mathcal{Y}$.
- The pairs (x, y) (observed data) are random samples of the random variables pair (X, Y) which has the distribution P_{XY} .

Definition 1.2 (Decomposition of P_{XY}).

We can decompose P_{XY} in either of the following two ways:

$$\begin{aligned} P_{XY} &= P_{X|Y} P_Y \\ P_{XY} &= P_{Y|X} P_X \end{aligned}$$

Which can be understood as two possible ways to generate the pairs (x, y) from the joint distribution P_{XY} .

- The first way is to generate a random label $y \sim P_Y$. Then, generate the feature vector corresponding to that label $x \sim P_{X|Y=y}$.
- The second way is to generate a random vector $x \sim P_X$. Then, generate the label corresponding to that feature vector $y \sim P_{Y|X=x}$.

Proposition 1.1: Law of total expectation

Given $\phi : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$. The **law of total expectation** states that:

$$\begin{aligned} \mathbb{E}_{XY} [\phi(X, Y)] &= \mathbb{E}_Y [\mathbb{E}_{X|Y} [\phi(X, Y)]] \\ &= \mathbb{E}_X [\mathbb{E}_{Y|X} [\phi(X, Y)]] \end{aligned}$$

Similar to how P_{XY} is decomposed, law of total expectation describes two way of taking the average value:

- Loop through the labels and take average over the feature vectors corresponding to each label.
- Loop through the feature vectors and take average over the labels corresponding to each vector.

Proof (Proposition 1.1).

We have:

$$\begin{aligned}
\mathbb{E}_{XY}[\phi(X, Y)] &= \int_{\mathcal{X}} \int_{\mathcal{Y}} \phi(x, y) P_{XY}(x, y) dy dx \\
&= \int_{\mathcal{X}} \int_{\mathcal{Y}} \phi(x, y) P_X(x) P_{Y|X}(y|x) dy dx \\
&= \int_{\mathcal{X}} P_X(x) \int_{\mathcal{Y}} \phi(x, y) P_{Y|X}(y|x) dy dx \\
&= \int_{\mathcal{X}} P_X(x) \mathbb{E}_{Y|X=x}[\phi(X, Y)] dx \\
&= \mathbb{E}_X[\mathbb{E}_{Y|X}[\phi(X, Y)]]
\end{aligned}$$

Applying the same technique, we have $\mathbb{E}_{XY}[\phi(X, Y)] = \mathbb{E}_Y[\mathbb{E}_{X|Y}[\phi(X, Y)]]$. \square .

Remark : Usually, the label space is discrete and finite, meaning $\mathcal{Y} = \{0, 1, 2, \dots, m\}$ for some $m < \infty$. Hence, the expectations over Y can be written as discrete sums:

$$\begin{aligned}
\mathbb{E}_{XY}[\phi(X, Y)] &= \mathbb{E}_Y[\mathbb{E}_{X|Y}[\phi(X, Y)]] = \sum_{y \in \mathcal{Y}} \mathbb{E}_{X|Y=y}[\phi(X, Y)] \\
&= \mathbb{E}_X[\mathbb{E}_{Y|X}[\phi(X, Y)]] = \mathbb{E}_X\left[\sum_{y \in \mathcal{Y}} \mathbb{E}_{Y=y|X}[\phi(X, Y)]\right]
\end{aligned}$$

Definition 1.3 (Hypothesis space (\mathcal{H})).

The hypothesis space is a collection (family) of classifiers $h : \mathcal{X} \rightarrow \mathcal{Y}$ that have some common properties:

$$\mathcal{H} = \left\{ h : \mathcal{X} \rightarrow \mathcal{Y} \mid \text{some common properties} \right\}$$

For example, let $\mathcal{X} = \mathbb{R}^d, \mathcal{Y} = (0, 1)$. In logistic regression, we assume the classifiers to be logit functions:

$$\mathcal{H}_{\text{logit}} = \left\{ h : \mathbb{R}^d \rightarrow (0, 1) \mid h(x) = \text{logit}(\beta x) = \frac{1}{1 + e^{-\beta x}}, \beta \in \mathbb{R}^{1 \times d} \right\}$$

Definition 1.4 (Learning algorithm (\mathcal{L}_n)).

To learn a classifier $h : \mathcal{X} \rightarrow \mathcal{Y}$, suppose that we have access to a training dataset of n data pairs $\{(X_k, Y_k)\}_{k=1}^n$ which are assumed to be **i.i.d sampled from** P_{XY} . The domain of the training data is then $(\mathcal{X} \times \mathcal{Y})^n$. A **learning algorithm**, denoted as \mathcal{L}_n is a function/procedure that derives a classifier $\hat{h}_n : \mathcal{X} \rightarrow \mathcal{Y}$ from the training data.

$$\begin{aligned}
\mathcal{L}_n &: (\mathcal{X} \times \mathcal{Y})^n \rightarrow \mathcal{H} \\
\hat{h}_n &= \mathcal{L}_n((X_1, Y_1), \dots, (X_n, Y_n))
\end{aligned}$$

1.2 Goal of classification

Definition 1.5 (Risk ($R(h)$)).

The **risk** of a classifier is defined as followed:

$$R(h) = P(h(X) \neq Y) = \mathbb{E}[\mathbf{1}_{\{h(X) \neq Y\}}]$$

Where (X, Y) are independent of the training data.

Definition 1.6 (Bayes Risk (R^*)).

The **Bayes risk** is the infimum of the risk taken over all $h : \mathcal{X} \rightarrow \mathcal{Y}$, not just for $h \in \mathcal{H}$:

$$R^* = \inf_{h: \mathcal{X} \rightarrow \mathcal{Y}} R(h)$$

Definition 1.7 (Consistency of learning algorithms).

A learning algorithm \mathcal{L}_n is called:

- **Weakly consistent** if $R(\hat{h}_n) \xrightarrow{P} R^*$:

$$\lim_{n \rightarrow \infty} P(R(\hat{h}_n) \leq r) = P(R^* \leq r), \quad \forall r \geq 0$$

- **Strongly consistent** if $R(\hat{h}_n) \xrightarrow{a.s.} R^*$:

$$P\left(\lim_{n \rightarrow \infty} |R(\hat{h}_n) - R^*| \geq \epsilon\right) = 0, \quad \forall \epsilon > 0$$

- **Universally weakly/strongly consistent** if \mathcal{L}_n is weakly/strongly consistent for all P_{XY} .
Meaning, consistency holds without any assumption about P_{XY} .

2 Bayes classifier

2.1 Properties of Bayes Risk

Overview : Recall that the Bayes classifier is the one with minimum risk and the corresponding risk is called the Bayes Risk. For $\mathcal{Y} = \{0, 1\}$ and defined:

$$\eta(x) = P(Y = 1|X = x)$$

Define the following classifier:

$$h^*(x) = \begin{cases} 1 & \text{if } \eta(x) \geq \frac{1}{2} \\ 0 & \text{otherwise} \end{cases}$$

Theorem 2.1: Properties of Bayes classifier

The following properties hold for the Bayes classifier with $\mathcal{Y} = \{0, 1\}$ (Binary classification):

- (i) $R(h^*) = \inf_{h: \mathcal{X} \rightarrow \mathcal{Y}} \{R(h)\} = R^*$.
- (ii) $\underbrace{R(h) - R^*}_{\text{Excess risk}} = 2\mathbb{E}_X \left[\left| \eta(x) - \frac{1}{2} \right| \mathbf{1}_{\{h(X) \neq h^*(X)\}} \right]$.
- (iii) $R^* = \mathbb{E} \left[\min(\eta(X), 1 - \eta(X)) \right]$.

Proof (Theorem 2.1). _____

Proving each point:

$$(i) \ R(h^*) = \inf_{h: \mathcal{X} \rightarrow \mathcal{Y}} \{R(h)\} = R^*.$$

For all $h: \mathcal{X} \rightarrow \mathcal{Y}$, we have:

$$\begin{aligned} R(h) &= \mathbb{E}_{XY} \left[\mathbf{1}_{\{h(X) \neq Y\}} \right] \\ &= \mathbb{E}_{x \sim X} \left[\mathbb{E}_{Y|X=x} \left[\mathbf{1}_{\{Y \neq h(x)\}} \right] \right] \\ &= \mathbb{E}_{x \sim X} \left[\sum_{y \in \{0,1\}} \mathbf{1}_{\{y \neq h(x)\}} \right] \\ &= \mathbb{E}_{x \sim X} \left[\eta(x) \mathbf{1}_{\{h(x)=0\}} + (1 - \eta(x)) \mathbf{1}_{\{h(x)=1\}} \right] \end{aligned}$$

Since the two events $\{h(x) = 1\}$ and $\{h(x) = 0\}$ are mutually exclusive, $R(h)$ is the smallest when we set $h(x) = 1$ when $\eta(x) \geq 1 - \eta(x) \implies \eta(x) \geq \frac{1}{2}$. Therefore, we have:

$$h^*(x) = \begin{cases} 1 & \text{if } \eta(x) \geq \frac{1}{2} \\ 0 & \text{otherwise} \end{cases}$$

$$(ii) \ \underbrace{R(h) - R^*}_{\text{Excess risk}} = 2\mathbb{E}_X \left[\left| \eta(x) - \frac{1}{2} \right| \mathbf{1}_{\{h(X) \neq h^*(X)\}} \right].$$

We have:

$$\begin{aligned}
R(h) - R^* &= \mathbb{E}_{x \sim X} \left[\mathbb{E}_{Y|X=x} \left[\mathbf{1}_{\{Y \neq h(x)\}} \right] \right] - \mathbb{E}_{x \sim X} \left[\mathbb{E}_{Y|X=x} \left[\mathbf{1}_{\{Y \neq h^*(x)\}} \right] \right] \\
&= \mathbb{E}_{x \sim X} \left[\sum_{y \in \{0,1\}} \mathbf{1}_{\{y \neq h(x)\}} P(Y = y|X = x) \right] - \mathbb{E}_{x \sim X} \left[\sum_{y \in \{0,1\}} \mathbf{1}_{\{y \neq h^*(x)\}} P(Y = y|X = x) \right] \\
&= \mathbb{E}_{x \sim X} \left[\eta(x) \left(\mathbf{1}_{\{h(x)=0\}} - \mathbf{1}_{\{h^*(x)=0\}} \right) + (1 - \eta(x)) \left(\mathbf{1}_{\{h(x)=1\}} - \mathbf{1}_{\{h^*(x)=1\}} \right) \right] \\
&= \mathbb{E}_{x \sim X} \left[\eta(x) \left(\mathbf{1}_{\{h(x) \neq h^*(x), h(x)=0\}} - \mathbf{1}_{\{h(x) \neq h^*(x), h(x)=1\}} \right) \right. \\
&\quad \left. + (1 - \eta(x)) \left(\mathbf{1}_{\{h(x) \neq h^*(x), h(x)=1\}} - \mathbf{1}_{\{h(x) \neq h^*(x), h(x)=0\}} \right) \right] \\
&= \mathbb{E}_{x \sim X} \left[(2\eta(x) - 1) \mathbf{1}_{\{h(x) \neq h^*(x), h(x)=0\}} + (1 - 2\eta(x)) \mathbf{1}_{\{h(x) \neq h^*(x), h(x)=1\}} \right] \\
&= \mathbb{E}_{x \sim X} \left[\left| 2\eta(x) - 1 \right| \mathbf{1}_{\{h(x) \neq h^*(x)\}} \right] \\
&= 2\mathbb{E}_X \left[\left| \eta(X) - \frac{1}{2} \right| \mathbf{1}_{\{h(X) \neq h^*(X)\}} \right]
\end{aligned}$$

$$(iii) \ R^* = \mathbb{E} \left[\min(\eta(X), 1 - \eta(x)) \right].$$

From (i) we have:

$$\begin{aligned}
R(h^*) &= \mathbb{E}_{x \sim X} \left[\eta(x) \mathbf{1}_{\{h^*(x)=0\}} + (1 - \eta(x)) \mathbf{1}_{\{h^*(x)=1\}} \right] \\
&= \mathbb{E}_X \left[\min(\eta(X), 1 - \eta(x)) \right]
\end{aligned}$$

□.

Theorem 2.2: Properties of Bayes classifier (Multi-class)

For multi-class classification with more than two labels : $\mathcal{Y} = \{1, 2, \dots, M\}$, the Bayes classifier is defined as followed:

$$h^*(x) = \arg \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\}$$

$$\text{Where : } \eta_y(x) = P(Y = y|X = x)$$

The following properties hold for the Bayes classifier with $\mathcal{Y} = \{1, 2, \dots, M\}$ (Multi-class classification):

- (i) **Bayes Risk** R^* :

$$R^* = \mathbb{E}_{x \sim X} \left[1 - \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} \right] = \mathbb{E}_{x \sim X} \left[\min_{y \in \mathcal{Y}} \overline{\eta}_y(x) \right]$$

- (ii) **Excess Risk** $R(h) - R^*$:

$$R(h) - R^* = \mathbb{E}_X \left[\left(\eta_{y_x^*}(x) - \eta_{y_x}(x) \right) \mathbf{1}_{\{h(x) \neq h^*(x)\}} \right]$$

Where $y_x = h(x)$ is the prediction made by an arbitrary classifier $h : \mathcal{X} \rightarrow \mathcal{Y}$ and $y_x^* = h^*(x)$ is the prediction made by the Bayes classifier.

Proof (Theorem 2.2).

(The proof of this theorem has been included in the solution of Exercise 2.1).

□.

2.2 Likelihood Ratio Test

Overview : Define $\pi_1 = P(Y = 1)$ and $\pi_0 = P(Y = 0)$ be the prior probabilities. Let $p_1(x) = P(X = x|Y = 1)$ and $p_0(x) = P(X = x|Y = 0)$ be the class-conditional densities. Note that we have:

$$\begin{aligned} \eta(x) &= P(Y = 1|X = x) \\ &= \frac{P(X = x|Y = 1)P(Y = 1)}{P(X = x|Y = 1)P(Y = 1) + P(X = x|Y = 0)P(Y = 0)} \\ &= \frac{\pi_1 p_1(x)}{\pi_1 p_1(x) + \pi_0 p_0(x)} \\ &= \frac{1}{1 + \frac{\pi_0 p_0(x)}{\pi_1 p_1(x)}} \end{aligned}$$

Hence, we have:

$$\begin{aligned} \eta(x) \geq \frac{1}{2} &\iff \frac{\pi_0 p_0(x)}{\pi_1 p_1(x)} \\ &\iff \frac{p_1(x)}{p_0(x)} \geq \frac{\pi_0}{\pi_1} \end{aligned}$$

Proposition 2.1: Likelihood ratio test

The Bayes classifier h^* can be re-defined as followed:

$$h^*(x) = \begin{cases} 1 & \text{if } \frac{p_1(x)}{p_0(x)} \geq \frac{\pi_0}{\pi_1} \\ 0 & \text{otherwise} \end{cases}$$

The fraction $\frac{p_1(x)}{p_0(x)}$ is called the **likelihood ratio**.

2.3 Plug-in classifier

Definition 2.1 (Plug-in classifier).

A **plug-in classifier** is based on an estimate of $\eta(x)$. This estimate is then plugged into the definition of the Bayes classifier. Suppose that $\widehat{\eta}_n$ is an estimate of η based on n training samples $\{(X_i, Y_i)\}_{i=1}^n$. We define \widehat{h}_n as:

$$\widehat{h}_n = \begin{cases} 1 & \text{if } \widehat{\eta}_n(x) \geq \frac{1}{2} \\ 0 & \text{otherwise} \end{cases}$$

Corollary 2.1: Excess risk of plug-in classifier

We have the following upper-bound for the excess risk of the plug-in classifier:

$$R(\widehat{h}_n) - R^* \leq 2\mathbb{E}_X \left[\left| \eta(X) - \widehat{\eta}_n(X) \right| \right]$$

Proof (Corollary 2.1).

From theorem 2.1, we have:

$$R(\widehat{h}_n) - R^* = 2\mathbb{E}_X \left[\left| \eta(X) - \frac{1}{2} \right| \mathbf{1}_{\{\widehat{h}_n(X) \neq h^*(X)\}} \right]$$

The indicator term will be non-zero in the above equality if one of the following cases occurs:

$$\begin{cases} \widehat{h}_n(X) = 1, h^*(X) = 0 \\ \widehat{h}_n(X) = 0, h^*(X) = 1 \end{cases} \implies \begin{cases} \widehat{\eta}_n(X) \geq \frac{1}{2}, \eta(X) < \frac{1}{2} \\ \widehat{\eta}_n(X) < \frac{1}{2}, \eta(X) \geq \frac{1}{2} \end{cases}$$

Case 1 : $\widehat{\eta}_n(X) \geq \frac{1}{2}, \eta(X) < \frac{1}{2}$

We have:

$$\begin{aligned} \eta(X) - \widehat{\eta}_n(X) &\leq \eta(X) - \frac{1}{2} \quad (\text{Both sides negative}) \\ \implies \left| \eta(X) - \widehat{\eta}_n(X) \right| &\geq \left| \eta(X) - \frac{1}{2} \right| \end{aligned}$$

Case 2 : $\widehat{\eta}_n(X) < \frac{1}{2}, \eta(X) \geq \frac{1}{2}$

We have:

$$\widehat{\eta}_n(X) - \eta(X) \geq \widehat{\eta}_n(X) - \frac{1}{2} \geq \eta(X) - \frac{1}{2} \quad (\text{All positive})$$

Therefore, we have:

$$\left| \eta(X) - \widehat{\eta}_n(X) \right| \geq \left| \eta(X) - \frac{1}{2} \right|$$

For both cases, we have the same $\left| \eta(X) - \widehat{\eta}_n(X) \right| \geq \left| \eta(X) - \frac{1}{2} \right|$ inequality. Therefore, we have:

$$R(\widehat{h}_n) - R^* \leq 2\mathbb{E}_X \left[\left| \eta(X) - \widehat{\eta}_n(X) \right| \right]$$

□.

2.4 End of chapter exercises

Exercise 2.1

Extend theorem 2.1 to the multi-class classification case where $\mathcal{Y} = \{1, 2, \dots, M\}$. In other words, prove theorem 2.2.

Solution (Exercise 2.1).

We re-define the Bayes classifier h^* as followed:

$$h^*(x) = \arg \max_{y \in \mathcal{Y}} \{\eta_y(x)\},$$

$$\eta_y(x) = P(Y = y | X = x)$$

We have:

$$\sum_{y \in \mathcal{Y}} \eta_y(x) = 1, \quad \forall x \in \mathcal{X}$$

(i) **Calculate Bayes risk R^***

For any classifier $h : \mathcal{X} \rightarrow \mathcal{Y}$, we have:

$$R(h) = \mathbb{E}_{x \sim X} \left[\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) \right]$$

Letting $\hat{y}_x = h(x)$ being h 's prediction for a given feature vector $x \in \mathcal{X}$, we have:

$$R(h) = \mathbb{E}_{x \sim X} \left[\sum_{y \in \mathcal{Y}; y \neq \hat{y}_x} \eta_y(x) \right] = \mathbb{E}_{x \sim X} \left[1 - \eta_{\hat{y}_x}(x) \right]$$

In order to minimize $R(h)$, we need $\eta_{\hat{y}_x}(x)$ to be maximized for all $x \in \mathcal{X}$. Hence, we have:

$$R^* = \mathbb{E}_{x \sim X} \left[1 - \max_{y \in \mathcal{Y}} \{\eta_y(x)\} \right]$$

Therefore, we have $h^*(x) = \arg \max_{y \in \mathcal{Y}} \{\eta_y(x)\}$ is the Bayes classifier and the Bayes risk $R^* = \mathbb{E}_{x \sim X} \left[1 - \max_{y \in \mathcal{Y}} \{\eta_y(x)\} \right]$.

(ii) **Calculate excess risk $R(h) - R^*$**

For any $h : \mathcal{X} \rightarrow \mathcal{Y}$, we have:

$$\begin{aligned} R(h) - R^* &= \mathbb{E}_{x \sim X} \left[\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) \right] - \mathbb{E}_{x \sim X} \left[1 - \max_{y \in \mathcal{Y}} \{\eta_y(x)\} \right] \\ &= \mathbb{E}_{x \sim X} \left[\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) + \max_{y \in \mathcal{Y}} \{\eta_y(x)\} - 1 \right] \end{aligned}$$

Denote $h^*(x) = y_x^*$ and $h(x) = y_x$. When $h(x) = h^*(x) = y_x^*$, we have:

$$\begin{aligned}
\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) + \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} &= \sum_{y \in \mathcal{Y}; y \neq y_x} \eta_y(x) + \eta_{y_x^*}(x) \\
&= \sum_{y \in \mathcal{Y}; y \neq y_x^*} \eta_y(x) + \eta_{y_x^*}(x) \\
&= \sum_{y \in \mathcal{Y}} \eta_y(x) = 1 \\
\Rightarrow \sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) + \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} - 1 &= 0
\end{aligned}$$

When $h(x) \neq h^*(x)$, we have:

$$\begin{aligned}
\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) + \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} - 1 &= \sum_{y \in \mathcal{Y}; y \neq y_x} \eta_y(x) + \eta_{y_x^*}(x) - 1 \\
&= 2\eta_{y_x^*}(x) - 1 + \sum_{y \in \mathcal{Y} \setminus \{y_x, y_x^*\}} \eta_y(x) \\
&= 2\eta_{y_x^*}(x) - \left(\eta_{y_x}(x) + \eta_{y_x^*}(x) \right) \\
&= \eta_{y_x^*}(x) - \eta_{y_x}(x).
\end{aligned}$$

Therefore, we can re-write the excess risk by multiplying the entire integrand with the indicator function $\mathbf{1}_{\{h(x) \neq h^*(x)\}}$ as followed:

$$R(h) - R^* = \mathbb{E}_{x \sim X} \left[\left(\eta_{y_x^*}(x) - \eta_{y_x}(x) \right) \mathbf{1}_{\{h(x) \neq h^*(x)\}} \right]$$

(iii) *Simpler form of Bayes risk*

From (i) we have:

$$R^* = \mathbb{E}_X \left[1 - \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} \right] = \mathbb{E}_X \left[\min_{y \in \mathcal{Y}} \left\{ \overline{\eta}_y(x) \right\} \right]$$

Where $\overline{\eta}_y(x) = P(Y \neq y | X = x)$.

□.

Exercise 2.2

Define the α -cost-sensitive risk of a classifier $h : \mathcal{X} \rightarrow \mathcal{Y}$ as followed:

$$R_\alpha(h) = \mathbb{E}_{XY} \left[(1 - \alpha) \mathbf{1}_{\{Y=1, h(X)=0\}} + \alpha \mathbf{1}_{\{Y=0, h(X)=1\}} \right]$$

Define the Bayes classifier and prove an analogue of theorem 2.1.

Solution (Exercise 2.2).

Using the law of total expectation, we have:

$$\begin{aligned}
R_\alpha(h) &= \mathbb{E}_{x \sim X} \left[\sum_{y \in \{0,1\}} \left[(1 - \alpha) \mathbf{1}_{\{y=1, h(x)=0\}} + \alpha \mathbf{1}_{\{y=0, h(x)=1\}} \right] P(Y = y | X = x) \right] \\
&= \mathbb{E}_{x \sim X} \left[(1 - \alpha) \eta(x) \mathbf{1}_{\{h(x)=0\}} + \alpha (1 - \eta(x)) \mathbf{1}_{\{h(x)=1\}} \right]
\end{aligned}$$

Since $\mathbf{1}_{\{h(x)=0\}}$ and $\mathbf{1}_{\{h(x)=1\}}$ are mutually exclusive, in order for $R_\alpha(h)$ to be minimize, we define the following Bayes classifier:

$$h^*(x) = \begin{cases} 1 & \text{if } \alpha(1 - \eta(x)) \leq (1 - \alpha)\eta(x) \\ 0 & \text{otherwise} \end{cases} = \begin{cases} 1 & \text{if } \eta(x) \geq \alpha \\ 0 & \text{otherwise} \end{cases}$$

We can also derive a likelihood-ratio test version of the Bayes classifier, we have:

$$\begin{aligned} \eta(x) \geq \alpha &\implies \frac{1}{1 + \frac{\pi_0 p_0(x)}{\pi_1 p_1(x)}} \geq \alpha \\ &\implies 1 + \frac{\pi_0 \cdot p_0(x)}{\pi_1 \cdot p_1(x)} \leq \frac{1}{\alpha} \\ &\implies \frac{p_1(x)}{p_0(x)} \geq \frac{\alpha}{1 - \alpha} \cdot \frac{\pi_0}{\pi_1} \end{aligned}$$

Hence, we can rewrite the Bayes classifier as followed:

$$h^*(x) = \begin{cases} 1 & \text{if } \frac{p_1(x)}{p_0(x)} \geq \frac{\alpha}{1 - \alpha} \cdot \frac{\pi_0}{\pi_1} \\ 0 & \text{otherwise} \end{cases}$$

(i) **Bayes Risk** R_α^*

We have:

$$\begin{aligned} R_\alpha^* &= R_\alpha(h^*) \\ &= \mathbb{E}_{x \sim X} \left[(1 - \alpha)\eta(x)\mathbf{1}_{\{h^*(x)=0\}} + \alpha(1 - \eta(x))\mathbf{1}_{\{h^*(x)=1\}} \right] \\ &= \mathbb{E}_X \left[\min(\alpha(1 - \eta(X)), (1 - \alpha)\eta(X)) \right] \end{aligned}$$

(ii) **Excess Risk** $R_\alpha(h) - R_\alpha^*$

For an arbitrary $h : \mathcal{X} \rightarrow \mathcal{Y}$, we have:

$$\begin{aligned} R_\alpha(h) - R_\alpha^* &= \mathbb{E}_{x \sim X} \left[(1 - \alpha)\eta(x) \left(\mathbf{1}_{\{h(x)=0\}} - \mathbf{1}_{\{h^*(x)=0\}} \right) + \alpha(1 - \eta(x)) \left(\mathbf{1}_{\{h(x)=1\}} - \mathbf{1}_{\{h^*(x)=1\}} \right) \right] \\ &= \mathbb{E}_{x \sim X} \left[(1 - \alpha)\eta(x) \left(\mathbf{1}_{\{h(x)=0, h^*(x)=1\}} - \mathbf{1}_{\{h(x)=1, h^*(x)=0\}} \right) \right. \\ &\quad \left. + \alpha(1 - \eta(x)) \left(\mathbf{1}_{\{h(x)=1, h^*(x)=0\}} - \mathbf{1}_{\{h(x)=0, h^*(x)=1\}} \right) \right] \\ &= \mathbb{E}_{x \sim X} \left[\mathbf{1}_{\{h(x)=0, h^*(x)=1\}} (\eta(x) - \alpha) + \mathbf{1}_{\{h(x)=1, h^*(x)=0\}} (\alpha - \eta(x)) \right] \\ &= \mathbb{E}_X \left[\left| \eta(X) - \alpha \right| \mathbf{1}_{\{h(X) \neq h^*(X)\}} \right] \end{aligned}$$

□.

3 Hoeffding's inequality

3.1 Markov's Inequality

Proposition 3.1: Markov's Inequality

Let U be a non-negative random variable on \mathbb{R} , then for all $t > 0$, we have:

$$P(U \geq t) \leq \frac{1}{t} \mathbb{E}[U]$$

Proof (Proposition 3.1). _____

We have:

$$\begin{aligned} tP(U \geq t) &= t\mathbb{E}[\mathbf{1}_{\{U \geq t\}}] \\ &= t \int_0^\infty \mathbf{1}_{\{x \geq t\}} f_U(x) dx \\ &= t \int_t^\infty f_U(x) dx \\ &\leq \int_t^\infty x f_U(x) dx \\ &\leq \int_0^\infty x f_U(x) dx = \mathbb{E}[U] \\ \implies P(U \geq t) &\leq \frac{1}{t} \mathbb{E}[U] \end{aligned}$$

□.

Corollary 3.1: Chebyshev's Inequality

Let Z be a random variable on \mathbb{R} with mean μ and variance σ^2 , we have:

$$P(|Z - \mu| \geq t) \leq \frac{\sigma^2}{t^2}$$

Proof (Corollary 3.1). _____

Using Markov's inequality, we have:

$$\begin{aligned} P(|Z - \mu| \geq t) &= P(|Z - \mu|^2 \geq t^2) \\ &\leq \frac{\mathbb{E}[|Z - \mu|^2]}{t^2} = \frac{\sigma^2}{t^2} \end{aligned}$$

□.

Corollary 3.2: Chernoff's bounding method

Let Z be a random variable on \mathbb{E} , for any $t > 0$, we have:

$$P(Z \geq t) \leq \inf_{s>0} e^{-st} M_Z(s)$$

Proof (Corollary 3.2). _____

We have:

$$\begin{aligned}
P(Z \geq t) &= P(sZ \geq st), \quad (t > 0) \\
&= P(e^{sZ} \geq e^{st}) \\
&\leq \frac{\mathbb{E}[e^{sZ}]}{e^{st}} = e^{-st} M_Z(s) \quad (\text{Markov's inequality})
\end{aligned}$$

Since the above inequality holds for all $s > 0$, we can just take the infimum to obtain the tightest bound. Hence, we have:

$$P(Z \geq t) \leq \inf_{s>0} e^{-st} M_Z(s)$$

□.

3.2 Hoeffding's Inequality

Before diving into Hoeffding's inequality, we need to go through the following lemma (whose proof will not be included) that will help us prove the Hoeffding's inequality:

Lemma 3.1: Hoeffding's lemma

Let V be a random variable on \mathbb{R} with $\mathbb{E}[V] = 0$ and suppose that $a \leq V \leq b$ with probability one. We have:

$$\mathbb{E}[e^{sV}] \leq \exp\left(\frac{s^2(b-a)^2}{8}\right)$$

Proof (Lemma 3.1). _____
(The proof for this lemma can be found here [4]).

□.

Theorem 3.1: Hoeffding's Inequality

Let Z_1, Z_2, \dots, Z_n be independent random variables on \mathbb{R} such that $a_i \leq Z_i \leq b_i$ with probability one for all $1 \leq i \leq n$. Let $S_n = \sum_{i=1}^n Z_i$. We have:

$$P\left(|S_n - \mathbb{E}[S_n]| \geq t\right) \leq 2 \exp\left(-\frac{2t^2}{\sum_{i=1}^n (b_i - a_i)^2}\right), \quad \forall t > 0$$

Proof (Theorem 3.1). _____

Using the Chernoff's bounds, we have:

$$\begin{aligned}
P\left(\left|S_n - \mathbb{E}[S_n]\right| \geq t\right) &\leq \inf_{s>0} e^{-st} M_{S_n - \mathbb{E}[S_n]}(s) \\
&= \inf_{s>0} e^{-st} \mathbb{E}\left[e^{s(S_n - \mathbb{E}[S_n])}\right] \\
&= \inf_{s>0} e^{-st} \mathbb{E}\left[\exp\left(s \sum_{i=1}^n (Z_i - \mathbb{E}[Z_i])\right)\right] \\
&= \inf_{s>0} e^{-st} \mathbb{E}\left[\prod_{i=1}^n \exp\left(s(Z_i - \mathbb{E}[Z_i])\right)\right] \\
&= \inf_{s>0} e^{-st} \prod_{i=1}^n \mathbb{E}\left[\exp\left(s(Z_i - \mathbb{E}[Z_i])\right)\right] \quad (\text{Since all } Z_i - \mathbb{E}[Z_i] \text{ are independent}) \\
&\leq \inf_{s>0} e^{-st} \prod_{i=1}^n \exp\left(\frac{s^2(b_i - a_i)^2}{8}\right) \quad (\text{By Hoeffding's lemma}) \\
&= \inf_{s>0} \exp\left(-st + \sum_{i=1}^n \frac{s^2(b_i - a_i)^2}{8}\right)
\end{aligned}$$

In order for the above to be minimized, we differentiate the term inside the exponential and set the derivative to 0 to find the optimal $s > 0$. We have:

$$-t + s \sum_{i=1}^n \frac{(b_i - a_i)^2}{4} = 0 \implies s = \frac{4t}{\sum_{i=1}^n (b_i - a_i)^2}$$

Letting $c = \sum_{i=1}^n (b_i - a_i)^2$, we now can derive the tightest Chernoff's bound as followed:

$$\begin{aligned}
P\left(\left|S_n - \mathbb{E}[S_n]\right| \geq t\right) &\leq \exp\left(-\frac{4t^2}{c} + \frac{16t^2}{c^2} \cdot \frac{c}{8}\right) = \exp\left(-\frac{2t^2}{c}\right) \\
&= \exp\left(-\frac{2t^2}{\sum_{i=1}^n (b_i - a_i)^2}\right)
\end{aligned}$$

□.

3.3 Convergence of Empirical Risk

Definition 3.1 (Empirical Risk (\widehat{R}_n)).

Suppose we are given training data $\{(X_i, Y_i)_{i=1}^n\}$ such that each pair $(X_i, Y_i) \sim P_{XY}$ are independently identically distributed. Let $h : \mathcal{X} \rightarrow \mathcal{Y}$ be a classifier. We define the **empirical risk** to be:

$$\widehat{R}_n(h) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}_{\{h(X_i) \neq Y_i\}}$$

Note that $\mathbb{E}[\widehat{R}_n(h)] = R(h)$ and $n\widehat{R}_n(h) \sim \text{Binomial}(n, R(h))$. In the following corollary of the Hoeffding's inequality, we will answer the question **how close the empirical risk is as an estimate of true risk** or **how fast the empirical risk converges to the true risk**.

Corollary 3.3: Convergence of Empirical Risk

Given training data $\{(X_i, Y_i)_{i=1}^n\}$ such that each pair $(X_i, Y_i) \sim P_{XY}$ are independently identically distributed. Let $h : \mathcal{X} \rightarrow \mathcal{Y}$ be a classifier, we have:

$$P\left(\left|\widehat{R}_n(h) - R(h)\right| \geq \epsilon\right) \leq 2e^{-2n\epsilon^2}, \quad \epsilon > 0$$

Proof (Corollary 3.3).

For all $1 \leq i \leq n$, we have $\mathbf{1}_{\{h(X_i) \neq Y_i\}} \in \{0, 1\}$. Hence, with probability one, $0 \leq \mathbf{1}_{\{h(X_i) \neq Y_i\}} \leq 1$ and $b_i = 1, a_i = 0$ for all $1 \leq i \leq n$.

Using the Hoeffding's inequality, we have:

$$\begin{aligned} P\left(\left|\widehat{R}_n(h) - R(h)\right| \geq \epsilon\right) &= P\left(\left|\widehat{R}_n(h) - \mathbb{E}[\widehat{R}_n(h)]\right| \geq \epsilon\right) \\ &= P\left(\left|n\widehat{R}_n(h) - \mathbb{E}[n\widehat{R}_n(h)]\right| \geq n\epsilon\right) \\ &\leq \exp\left(-\frac{2n^2\epsilon^2}{\sum_{i=1}^n (b_i - a_i)^2}\right) \quad (\text{Hoeffding's inequality}) \\ &= e^{-2n\epsilon^2} \end{aligned}$$

□.

3.4 KL-divergence & Hypothesis Testing

Set-up (Hypothesis Testing) : Suppose that we have $\mathcal{Y} = \{0, 1\}$ and P_{XY} is a distribution on $\mathcal{X} \times \mathcal{Y}$. Let's assume that:

- The prior probabilities π_y are equal.
- The supports of likelihoods p_0, p_1 are the same.
- $0 < \alpha \leq p_y(x) \leq \beta < \infty$ for all $x \in \mathcal{X}$ such that $p_y(x) > 0$ and for all $y \in \{0, 1\}$.

Now suppose $X_1, \dots, X_n \sim p_y$ are independently identically distributed where $y \in \{0, 1\}$ is unknown. Can we guess y and how good our guess would be?

Proposition 3.2: KL-divergence hypothesis testing

From the above settings, the optimal classifier is given by the likelihood ratio test:

$$\widehat{h}_n(x) = \begin{cases} 1 & \text{if } \frac{\prod_{i=1}^n p_1(x_i)}{\prod_{i=1}^n p_0(x_i)} \geq \frac{\pi_0}{\pi_1} \quad (= 1) \\ 0 & \text{otherwise} \end{cases}$$

Where $x = (x_1, \dots, x_n)$ is an observation of the random vector $X = (X_1, \dots, X_n)$. Define the class-specific risk $R_y(h)$ be the risk of misclassification when the true label is $Y = y$:

$$R_y(h) = P(h(X) \neq Y | Y = y)$$

Then, we have:

$$R_0(\widehat{h}_n) \leq e^{-2nD(p_0||p_1)^2/c}, \text{ where } c = 4(\log \beta - \log \alpha)^2$$

Where $D(p_0||p_1)$ is the *KL*-divergence of p_1 from p_0 . We can prove a similar exponentially decaying bound for $R_1(\widehat{h}_n)$.

Proof.

Proposition 3.2 We can rewrite the optimal classifier as:

$$\widehat{h}_n(X) = \begin{cases} 1 & \text{if } \widehat{S}_n(X_1, \dots, X_n) \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

Where we have:

$$\begin{aligned} \widehat{S}_n(X_1, \dots, X_n) &= \log \frac{\prod_{i=1}^n p_1(X_i)}{\prod_{i=1}^n p_0(X_i)} \\ &= \sum_{i=1}^n \log \frac{p_1(X_i)}{p_0(X_i)} \\ &= \sum_{i=1}^n Z_i \quad \left(\text{Letting } Z_i = \log \frac{p_1(X_i)}{p_0(X_i)} \right) \end{aligned}$$

Since the likelihoods are bounded, we have:

$$a_i = \log \frac{\alpha}{\beta} \leq Z_i \leq \log \frac{\beta}{\alpha} = b_i, \quad 1 \leq i \leq n$$

Now, we have:

$$\begin{aligned} R_0(\widehat{h}_n) &= P(h(X) \neq Y | Y = 0) \\ &= P(\widehat{S}_n \geq 0 | Y = 0) \\ &= P(\widehat{S}_n - \mathbb{E}[S_n | Y = 0] \geq -\mathbb{E}[S_n | Y = 0] | Y = 0) \end{aligned}$$

To calculate the conditional expectation $\mathbb{E}[S_n | Y = 0]$, we have:

$$\begin{aligned} \mathbb{E}[S_n | Y = 0] &= n\mathbb{E}[Z_1 | Y = 0] \\ &= n \int \log \frac{p_1(x)}{p_0(x)} p_0(x) dx \\ &= -n \int \log \frac{p_0(x)}{p_1(x)} p_0(x) dx = -nD(p_0||p_1) \end{aligned}$$

Therefore, we have:

$$\begin{aligned} R_0(\widehat{h}_n) &= P(\widehat{S}_n - \mathbb{E}[S_n|Y=0] \geq nD(p_0||p_1)|Y=0) \\ &\leq \exp\left(-\frac{2n^2D(p_0||p_1)^2}{\sum_{i=1}^n(b_i - a_i)^2}\right) \quad (\text{Hoeffding's inequality}) \end{aligned}$$

For every $1 \leq i \leq n$, we have:

$$\begin{aligned} b_i - a_i &= \log \frac{\beta}{\alpha} - \log \frac{\alpha}{\beta} \\ &= \log \frac{\beta^2}{\alpha^2} = 2 \log \frac{\beta}{\alpha} = 2(\log \beta - \log \alpha) \\ \implies \sum_{i=1}^n (b_i - a_i)^2 &= 4n(\log \beta - \log \alpha)^2 \end{aligned}$$

Finally, we have:

$$R_0(\widehat{h}_n) \leq \exp\left(-\frac{2nD(p_0||p_1)^2}{4(\log \beta - \log \alpha)^2}\right)$$

Similarly, for $R_1(\widehat{h}_n)$, we have:

$$R_1(\widehat{h}_n) \leq \exp\left(-\frac{2nD(p_1||p_0)^2}{4(\log \beta - \log \alpha)^2}\right)$$

□.

3.5 End of chapter exercises

Exercise 3.1

- (i) Apply Chernoff's bounding method to obtain an exponential bound on the tail probability $P(Z \geq t)$ for a Gaussian random variable $Z \sim \mathcal{N}(\mu, \sigma^2)$.
- (ii) Appealing to the central limit theorem, use part (i) to give an approximate bound on the binomial tail. This should not only match the exponential decay given by Hoeffding's inequality, but also reveal the dependence on the variance of the binomial.

Solution (Exercise 3.1). _____

(i) **Chernoff's bounds for $Z \sim \mathcal{N}(\mu, \sigma^2)$**

Using the Chernoff's bounding method, we have:

$$\begin{aligned} P(Z \geq t) &\leq \inf_{s>0} e^{-st} M_Z(s) \\ &= \inf_{s>0} \exp \left(-st + \mu s + \frac{1}{2} \sigma^2 s^2 \right) \end{aligned}$$

The above bound is the tightest when the derivative of the term inside the exponential equals zero. Hence, we have:

$$-t + \mu + s\sigma^2 = 0 \implies s = \frac{t - \mu}{\sigma^2}$$

From the above, we have the tightest Chernoff's bound as followed:

$$P(Z \geq t) \leq \exp \left(-\frac{(t - \mu)^2}{\sigma^2} + \frac{(t - \mu)^2}{2\sigma^2} \right) = \exp \left(-\frac{(t - \mu)^2}{2\sigma^2} \right)$$

(ii) **Binomial tail upper bound**

Let S_n be the binomial random variable such that:

$$S_n = \sum_{i=1}^n X_i, \quad X_i \sim \text{Bernoulli}(p)$$

For a positive $\epsilon > 0$, we want to know the upper tail bound $P(S_n - \mathbb{E}[S_n] \geq \epsilon)$. Letting $\bar{X} = \frac{1}{n} S_n$, we have:

$$\begin{aligned} P(S_n - \mathbb{E}[S_n] \geq \epsilon) &= P \left(\bar{X} - \frac{\mathbb{E}[S_n]}{n} \geq \frac{\epsilon}{n} \right) \\ &= P \left(\bar{X} - p \geq \frac{\epsilon}{n} \right) \\ &= P \left(\frac{\bar{X} - p}{\sqrt{pq}/\sqrt{n}} \geq \frac{\epsilon}{\sqrt{npq}} \right), \quad (q = 1 - p) \end{aligned}$$

By the Central Limit Theorem, we have:

$$\frac{\bar{X} - p}{\sqrt{pq}/\sqrt{n}} \xrightarrow{d} \mathcal{N}(0, 1)$$

Hence, as $n \rightarrow \infty$, the upper tail bound would be:

$$\begin{aligned} P(S_n - \mathbb{E}[S_n] \geq \epsilon) &= P\left(\frac{\bar{X} - p}{\sqrt{pq}/\sqrt{n}} \geq \frac{\epsilon}{\sqrt{npq}}\right) \\ &\leq \exp\left(-\frac{\epsilon^2}{2npq}\right) = \exp\left(-\frac{\epsilon^2}{2\text{Var}(S_n)}\right) \end{aligned}$$

Double-check the bound with Hoeffding's inequality, we have:

$$P(S_n - \mathbb{E}[S_n] \geq \epsilon) \leq \exp\left(-\frac{2\epsilon^2}{n}\right)$$

□.

Exercise 3.2

Can you remove the assumption in $0 < \alpha \leq p_y(x)$? Consider other restrictions on p_y , other concentration inequalities, or other f -divergences.

Solution (Exercise 3.2).

When we remove the assumption that $0 < \alpha \leq p_y(x)$, the class-conditional densities are not bounded below. Hence, we have:

$$\exp\left(-\frac{2nD(p_1||p_0)^2}{4(\log \beta - \log \alpha)^2}\right) \rightarrow 1 \text{ when } \alpha \rightarrow 0$$

In other words, the bound is no longer meaningful. We can instead use the Chernoff bounding method:

$$\begin{aligned} R_0(\widehat{h}_n) &= P(S_n \geq 0 | Y = 0) \\ &\leq \inf_{s>0} \prod_{i=1}^n \mathbb{E}_{q_0} \left[e^{sZ_i} \right] \\ &= \inf_{s>0} \prod_{i=1}^n \mathbb{E}_{q_0} \left[\exp \left(s \log \frac{p_1(X_i)}{p_0(X_i)} \right) \right] \\ &= \inf_{s>0} \prod_{i=1}^n \mathbb{E}_{q_0} \left[\frac{p_1(X_i)^s}{p_0(X_i)^s} \right] \end{aligned}$$

Taking logarithm from both sides, we have:

$$\begin{aligned} \log R_0(\widehat{h}_n) &\leq \inf_{s>0} \sum_{i=1}^n \log \mathbb{E}_{q_0} \left[\frac{p_1(X_i)^s}{p_0(X_i)^s} \right] \\ &= \inf_{s>0} \sum_{i=1}^n (s-1) R_s(p_1||p_0) \\ &= \inf_{s>0} n(s-1) R_s(p_1||p_0) \end{aligned}$$

Where $R_s(p_1||p_0)$ is the Renyi divergence [5].

□.

4 Empirical Risk Minimization

4.1 Uniform Deviation Bounds

Definition 4.1 (Empirical Risk Minimization (\widehat{h}_n)).

Let $\{(X_i, Y_i)\}_{i=1}^n$ be independently identically distributed random variables sampled from P_{XY} . Let $\mathcal{H} \subset \{0, 1\}^{\mathcal{X}}$ be a set of classifiers. **Empirical Risk Minimization** is a learning algorithm such that:

$$\widehat{h}_n = \arg \min_{h \in \mathcal{H}} \widehat{R}_n(h)$$

Where \widehat{R}_n is the empirical risk and \widehat{h}_n is called the **Empirical Risk Minimizer**. An important question is how close \widehat{R}_n is to $R_{\mathcal{H}}^* = \inf_{h \in \mathcal{H}} R(h)$.

Overview (Uniform Deviation Bounds) : Previously, we proved the following bound using the Hoeffding's inequality:

$$P\left(\left|\widehat{R}_n(h) - R(h)\right| \geq \epsilon\right) \leq \delta$$

Where $\delta = 2e^{-2n\epsilon^2}$. However, since we do not know \widehat{h}_n (the specific function in \mathcal{H} that minimizes the empirical risk), we look for a bound that is guaranteed to apply for all $h \in \mathcal{H}$. This is called the Uniform Deviation Bound.

Definition 4.2 (Uniform Deviation Bounds (UDB)).

Given a set of classifiers $\mathcal{H} \subset \{0, 1\}^{\mathcal{X}}$, $\epsilon > 0$, the **Uniform Deviation Bounds** is the probability that for at least one $h \in \mathcal{H}$, the empirical risk deviates away from the true risk by ϵ and has the following form:

$$P\left(\sup_{h \in \mathcal{H}} \left|\widehat{R}_n(h) - R(h)\right| \leq \epsilon\right) \geq 1 - \delta$$

$$\text{Or : } P\left(\sup_{h \in \mathcal{H}} \left|\widehat{R}_n(h) - R(h)\right| \geq \epsilon\right) \leq \delta$$

The above bounds have the following interpretations:

- The probability that the deviation from the true risk is at most ϵ for all functions in \mathcal{H} is at least $1 - \delta$.
- The probability that there exists at least a function in \mathcal{H} whose deviation from the true risk is at least ϵ is at most δ .

Basically, we want to **bound the probability that some function deviates too far from the true risk**.

Proposition 4.1: Uniform Deviation Bounds for finite \mathcal{H}

Assume that $|\mathcal{H}| < \infty$. We have:

$$P\left(\sup_{h \in \mathcal{H}} \left|\widehat{R}_n(h) - R(h)\right| \geq \epsilon\right) \leq 2|\mathcal{H}|e^{-2n\epsilon^2}$$

Proof (Proposition 4.1).

For $h \in \mathcal{H}$, define the following event:

$$\Omega_\epsilon(h) = \left\{ \left| \widehat{R}_n(h) - R(h) \right| \geq \epsilon \right\}$$

Which is the event that the function h deviates away from the true risk by $\epsilon > 0$. Now, define the following event:

$$\Omega_\epsilon(\mathcal{H}) = \bigcup_{h \in \mathcal{H}} \Omega_\epsilon(h)$$

Which is the event that at least one $h \in \mathcal{H}$ deviates away from the true risk by $\epsilon > 0$. We have:

$$\begin{aligned} P\left(\sup_{h \in \mathcal{H}} \left| \widehat{R}_n(h) - R(h) \right| \geq \epsilon\right) &= P(\Omega_\epsilon(\mathcal{H})) \\ &= P\left(\bigcup_{h \in \mathcal{H}} \Omega_\epsilon(h)\right) \\ &\leq \sum_{h \in \mathcal{H}} P(\Omega_\epsilon(h)) \\ &\leq \sum_{h \in \mathcal{H}} 2e^{-2n\epsilon^2} = 2|\mathcal{H}|e^{-2n\epsilon^2} \end{aligned}$$

□.

Proposition 4.2: (Probabilistic) Bound on Excess Risk of \widehat{h}_n

Suppose that \mathcal{H} satisfies:

$$P\left(\sup_{h \in \mathcal{H}} \left| \widehat{R}_n(h) - R(h) \right| \geq \epsilon\right) \leq \delta$$

Then, with probability of at least $1 - \delta$, we have the following **upper bound on the Excess Risk of the Empirical Risk Minimizer**:

$$R(\widehat{h}_n) - R_{\mathcal{H}}^* \leq 2\epsilon$$

In other words, **with probability $1 - \delta$, the empirical risk minimizer deviates from the true risk minimizer by at most 2ϵ .**

Proof (Proposition 4.2).

We have:

$$P\left(\sup_{h \in \mathcal{H}} \left| \widehat{R}_n(h) - R(h) \right| \geq \epsilon\right) \leq \delta \implies P\left(\sup_{h \in \mathcal{H}} \left| \widehat{R}_n(h) - R(h) \right| \leq \epsilon\right) \geq 1 - \delta$$

Hence, with probability $1 - \delta$, for all $h \in \mathcal{H}$, we have:

$$\begin{aligned} \left| \widehat{R}_n(h) - R(h) \right| \leq \epsilon &\implies -\epsilon \leq \widehat{R}_n(h) - R(h) \leq \epsilon \\ &\implies \begin{cases} \widehat{R}_n(h) &\leq R(h) + \epsilon \\ R(h) &\leq \widehat{R}_n(h) + \epsilon \end{cases} \end{aligned}$$

Therefore:

$$\begin{aligned}
R(\widehat{h}_n) &\leq \widehat{R}_n(\widehat{h}_n) + \epsilon \\
&\leq \widehat{R}_n(h) + \epsilon \quad (\text{Since } \widehat{h}_n \text{ minimizes the Empirical Risk}) \\
&\leq (R(h) + \epsilon) + \epsilon = R(h) + 2\epsilon
\end{aligned}$$

Since $h \in \mathcal{H}$ is an arbitrary choice, we take the infimum over \mathcal{H} to get the tightest bound. We have:

$$\begin{aligned}
R(\widehat{h}_n) &\leq \inf_{h \in \mathcal{H}} R(h) + 2\epsilon \\
&= R_{\mathcal{H}}^* + 2\epsilon
\end{aligned}$$

□.

Remark : We can express the above proposition verbally as "If the UDB is at most δ , then with probability $1 - \delta$, the Excess Risk of the Empirical Risk Minimizer is at most 2ϵ ".

Remark : Note that the above proof assumes that *there exists an empirical risk minimizer*. This is not guaranteed when $|\mathcal{H}|$ is infinite.

Proposition 4.3: (Non-probabilistic) Bound on Excess Risk of \widehat{h}_n

We have the following inequality:

$$R(\widehat{h}_n) - R_{\mathcal{H}}^* \leq 2 \sup_{h \in \mathcal{H}} |\widehat{R}_n(h) - R(h)|$$

Proof (Proposition 4.3).

Let $h_{\mathcal{H}}^* = \arg \min_{h \in \mathcal{H}} R(h)$. We have:

$$R(\widehat{h}_n) - R_{\mathcal{H}}^* \leq |R(\widehat{h}_n) - \widehat{R}_n(\widehat{h}_n)| + \widehat{R}_n(\widehat{h}_n) - \widehat{R}_n(h_{\mathcal{H}}^*) + |\widehat{R}_n(h_{\mathcal{H}}^*) - R_{\mathcal{H}}^*|$$

Since \widehat{h}_n is the Empirical Risk Minimizer, we have $\widehat{R}_n(\widehat{h}_n) - \widehat{R}_n(h_{\mathcal{H}}^*) \leq 0$. Hence:

$$\begin{aligned}
R(\widehat{h}_n) - R_{\mathcal{H}}^* &\leq |R(\widehat{h}_n) - \widehat{R}_n(\widehat{h}_n)| + |\widehat{R}_n(h_{\mathcal{H}}^*) - R_{\mathcal{H}}^*| \\
&\leq 2 \sup_{h \in \mathcal{H}} |\widehat{R}_n(h) - R(h)|
\end{aligned}$$

□.

Corollary 4.1: Excess Risk of \widehat{h}_n - $\delta \rightarrow \epsilon$ relation

This is a Corollary for both proposition 4.2 and proposition 4.3. If \mathcal{H} is finite, then:

$$P\left(R(\widehat{h}_n) - R_{\mathcal{H}}^* \geq \epsilon\right) \leq \underbrace{2|\mathcal{H}|e^{-n\epsilon^2/2}}_{\delta}$$

Equivalently, with probability of at least $1 - \delta$, we have:

$$R(\widehat{h}_n) \leq R_{\mathcal{H}}^* + \sqrt{\frac{2}{n} \left(\log |\mathcal{H}| - \log \frac{\delta}{2} \right)}$$

Proof (Corollary 4.1).

By proposition 4.3, we have:

$$\begin{aligned} P\left(R(\widehat{h}_n) - R_{\mathcal{H}}^* \geq \epsilon\right) &\leq P\left(2 \sup_{h \in \mathcal{H}} \left|\widehat{R}_n(h) - R(h)\right| \geq \epsilon\right) \\ &= P\left(\sup_{h \in \mathcal{H}} \left|\widehat{R}_n(h) - R(h)\right| \geq \frac{\epsilon}{2}\right) \\ &\leq 2|\mathcal{H}| \exp\left(-\frac{n\epsilon^2}{2}\right) \end{aligned}$$

Now, let:

$$\delta = 2|\mathcal{H}| \exp\left(-\frac{n\epsilon^2}{2}\right) \implies \epsilon = \sqrt{\frac{2}{n} \left(\log |\mathcal{H}| - \log \frac{\delta}{2}\right)}$$

By proposition 4.2, with at least probability $1 - \delta$, we have:

$$R(\widehat{h}_n) \leq R_{\mathcal{H}}^* + \epsilon = R_{\mathcal{H}}^* + \sqrt{\frac{2}{n} \left(\log |\mathcal{H}| - \log \frac{\delta}{2}\right)}$$

□.

4.2 PAC Learning & Sample Complexity

Definition 4.3 (PAC & Sample Complexity $(N(\epsilon, \delta))$).

We say that \widehat{h}_n is a (ϵ, δ) -**learning algorithm** for \mathcal{H} if there exists a function $N(\epsilon, \delta)$ such that:

$$\forall \epsilon, \delta > 0 : n \geq N(\epsilon, \delta) \implies P\left(R(\widehat{h}_n) - R_{\mathcal{H}}^* \geq \epsilon\right) \leq \delta$$

Where we have:

- $N(\epsilon, \delta)$ is called the **Sample Complexity**.
- \mathcal{H} is called **Uniformly Learnable**.
- \widehat{h}_n is called **Probably Approximately Correct (PAC)**.

Remark : By corollary 4.1, we have $\delta = 2|\mathcal{H}| \exp\left(-\frac{n\epsilon^2}{2}\right)$. Solving for n , we have:

$$N(\epsilon, \delta) = \frac{2}{\epsilon^2} \left(\log |\mathcal{H}| - \log \frac{\delta}{2}\right)$$

4.3 Zero-error case

In the following proposition, we can obtain a tighter bound for the zero empirical risk case. However, it is not particularly useful in many cases.

Proposition 4.4: Zero-error case bound

If $\widehat{R}_n(\widehat{h}_n) = 0$ and $|\mathcal{H}| < \infty$, we have:

$$P\left(\exists h \in \mathcal{H} : \widehat{R}_n(h) = 0, R(h) \geq \epsilon\right) \leq \underbrace{|\mathcal{H}|e^{-n\epsilon}}_{\delta}$$

Meaning, with probability of at least $1 - \delta$, if $\widehat{R}_n(h) = 0$ then $R(h) \leq \frac{1}{n}(\log |\mathcal{H}| - \log \delta)$.

Proof (Proposition 4.4).

Let $\Omega_0(h) = \{\widehat{R}_n(h) = 0\}$ and define the event Ω_ϵ as:

$$\Omega_\epsilon = \bigcup_{h \in \mathcal{H}; R(h) \geq \epsilon} \Omega_0(h) = \left\{ \exists h \in \mathcal{H} : \widehat{R}_n(h) = 0, R(h) \geq \epsilon \right\}$$

For any $h \in \mathcal{H}$ such that $R(h) \geq \epsilon$, we have:

$$\begin{aligned} P(\Omega_0(h)) &= P\left(\frac{1}{n} \sum_{i=1}^n \mathbf{1}_{\{h(X_i) \neq Y_i\}} = 0\right) \\ &= P\left(\sum_{i=1}^n \mathbf{1}_{\{h(X_i) \neq Y_i\}} = 0\right) \\ &= P\left(\bigcup_{i=1}^n \{h(X_i) = Y_i\}\right) \\ &= \prod_{i=1}^n P(h(X_i) = Y_i) \quad (\text{Since all } (X_i, Y_i) \text{ pairs are independent}) \end{aligned}$$

Each $\mathbf{1}_{\{h(X_i) \neq Y_i\}}$ is a Bernoulli variable with hit probability $p_i = 1 - \mathbb{E}[h(X_i) \neq Y_i] = 1 - R(h)$. Hence, we have:

$$\begin{aligned} P(\Omega_0(h)) &= \prod_{i=1}^n P(h(X_i) = Y_i) \\ &= (1 - R(h))^n \\ &\leq (1 - \epsilon)^n \end{aligned}$$

Using the inequality $\log(1 - \epsilon) \leq -\epsilon$, we have:

$$\begin{aligned} P(\Omega_0(h)) &\leq (1 - \epsilon)^n = e^{n \log(1 - \epsilon)} \\ &\leq e^{-n\epsilon} \end{aligned}$$

Finally, we have:

$$\begin{aligned} P(\Omega_\epsilon) &= P\left(\bigcup_{h \in \mathcal{H}; R(h) \geq \epsilon} \Omega_0(h)\right) \\ &\leq \sum_{h \in \mathcal{H}; R(h) \geq \epsilon} P(\Omega_0(h)) \\ &\leq \sum_{h \in \mathcal{H}; R(h) \geq \epsilon} e^{-n\epsilon} \\ &\leq |\mathcal{H}|e^{-n\epsilon} \end{aligned}$$

□.

Remark : Note that the bound obtained in proposition 4.4 is NOT the Uniform Deviation Bound (UDB) because we have:

$$\left\{ \sup_{h \in \mathcal{H}} \left| \widehat{R}_n(h) - R(h) \right| \geq \epsilon \right\} = \left\{ \exists h \in \mathcal{H} : \left| \widehat{R}_n(h) - R(h) \right| \geq \epsilon \right\}$$

Therefore, we have:

$$\left\{ \exists h \in \mathcal{H} : \widehat{R}_n(h) = 0, R(h) \geq \epsilon \right\} \subseteq \left\{ \sup_{h \in \mathcal{H}} \left| \widehat{R}_n(h) - R(h) \right| \geq \epsilon \right\}$$

Remark : This is trivial improvement. However, define the following subset of \mathcal{H} :

$$H_\epsilon^+ = \left\{ h \in \mathcal{H} : R(h) \geq \epsilon \right\}$$

We can improve the bound in proposition 4.4 as followed:

$$P(\Omega_\epsilon) \leq |H_\epsilon^+| e^{-n\epsilon}$$

4.4 End of chapter exercises

Exercise 4.1

The probability of error is not the only performance measure for binary classification. Indeed, the probability of error depends on the prior probability of the class label Y , and it may be that the frequency of the classes changes from training to testing data. In such cases, it is desirable to have a performance measure that does not require knowledge of the prior class probability. Let P_y be the class conditional distribution of class $y \in \{0, 1\}$. Define $R_y(h) = P_y(h(X) \neq y)$. Also let $\alpha \in (0, 1)$. For $\mathcal{H} \subset \{0, 1\}^{\mathcal{X}}$, define:

$$R_{\mathcal{H},1}^* = \inf_{h \in \mathcal{H}} R_1(h) \\ \text{s.t. } R_0(h) \leq \alpha$$

In this problem you will investigate a discrimination rule that is probably approximately correct with respect to the above criterion, which is sometimes called the Neyman-Pearson criterion based on connections to the Neyman-Pearson lemma in hypothesis testing. Suppose we observe $X_1^y, X_2^y, \dots, X_{n_y}^y \sim P_y$ for $y \in \{0, 1\}$. Define the empirical errors:

$$\widehat{R}_y(h) = \frac{1}{n_y} \sum_{i=1}^{n_y} \mathbf{1}_{\{h(X_i^y) \neq y\}}$$

Fix $\epsilon > 0$ and consider the discrimination rule:

$$\widehat{h}_n = \arg \min_{h \in \mathcal{H}} \widehat{R}_1(h) \\ \text{s.t. } \widehat{R}_0(h) \leq \alpha + \frac{\epsilon}{2}$$

Suppose \mathcal{H} is finite. Show that with high probability:

$$R_0(\widehat{h}_n) \leq \alpha + \epsilon \text{ and } R_1(\widehat{h}_n) \leq R_{\mathcal{H},1}^* + \epsilon$$

A Related topics

A.1 Neyman-Pearson Lemma

A.1.1 Type I & Type II errors

Overview : In a hypothesis test, we are interested in testing a given null hypothesis H_0 against some alternative hypothesis H_1 . Hence, we define some rejection region $\mathcal{R} \subset \mathbb{R}$ such that:

$$x \in \mathcal{R} \implies \text{reject } H_0$$

Equivalently, denote that $\bar{\mathcal{R}}$ is the acceptance region. We define the following conditional probability densities:

- $f_1(x)$: Density given that H_1 is true.
- $f_0(x)$: Density given that H_0 is true.

Definition A.1 (Type I & Type II errors).

In hypothesis testing, we define the type I error as the probability that we falsely reject the null hypothesis given that the null hypothesis is true. On the other hands, type II error is the probability that we falsely accept the null hypothesis given that the hypothesis is not true:

$$\alpha = P(\mathcal{R}|H_0) = \int_{\mathcal{R}} f_0(x)dx$$
$$\beta = P(\bar{\mathcal{R}}|H_1) = 1 - \int_{\mathcal{R}} f_1(x)dx$$

There is a trade-off between type I and type II errors as illustrated in the figure below:




Figure 1: Trade-off between type I and type II errors

Definition A.2 (Power of hypothesis test).

Given a hypothesis test used to test a null hypothesis H_0 against an alternative hypothesis H_1 . The probability:

$$P_{H_1}(\mathcal{R}) = \int_{\mathcal{R}} f_1(x)dx = 1 - \beta$$

Which denotes the probability that we correctly reject the null hypothesis given that H_1 is true is called the **Power** of the hypothesis test. Later on we will see that using **Neyman-Pearson Lemma**, we can prove any hypothesis test has the power of at most the likelihood ratio test's power.

A.1.2 Neyman-Pearson Lemma

B List of Definitions

1.1	Definition (Classifier (h))	2
1.2	Definition (Decomposition of P_{XY})	2
1.3	Definition (Hypothesis space (\mathcal{H}))	3
1.4	Definition (Learning algorithm (\mathcal{L}_n))	3
1.5	Definition (Risk ($R(h)$))	4
1.6	Definition (Bayes Risk (R^*))	4
1.7	Definition (Consistency of learning algorithms)	4
2.1	Definition (Plug-in classifier)	8
3.1	Definition (Empirical Risk (\widehat{R}_n))	15
4.1	Definition (Empirical Risk Minimization (\widehat{h}_n))	21
4.2	Definition (Uniform Deviation Bounds (UDB))	21
4.3	Definition (PAC & Sample Complexity ($N(\epsilon, \delta)$))	24
A.1	Definition (Type I & Type II errors)	28
A.2	Definition (Power of hypothesis test)	29

C Important Theorems

2.1	Properties of Bayes classifier	5
2.2	Properties of Bayes classifier (Multi-class)	7
3.1	Hoeffding's Inequality	14

D Important Corollaries

2.1	Excess risk of plug-in classifier	8
3.1	Chebyshev's Inequality	13
3.2	Chernoff's bounding method	13
3.3	Convergence of Empirical Risk	16
4.1	Excess Risk of \widehat{h}_n - $\delta \rightarrow \epsilon$ relation	23

E Important Propositions

1.1	Law of total expectation	2
2.1	Likelihood ratio test	8
3.1	Markov's Inequality	13
3.2	KL-divergence hypothesis testing	17
4.1	Uniform Deviation Bounds for finite \mathcal{H}	21
4.2	(Probabilistic) Bound on Excess Risk of \widehat{h}_n	22
4.3	(Non-probabilistic) Bound on Excess Risk of \widehat{h}_n	23
4.4	Zero-error case bound	25

F References

References

- [1] Rick Durrett. *Probability: Theory and Examples*. 4th. USA: Cambridge University Press, 2010. ISBN: 0521765390.
- [2] Zhou Fan. *Statistics 200: Introduction to Statistical Inference - Lecture 6*. <https://web.stanford.edu/class/archive/stats/stats200/stats200.1172/Lecture06.pdf>. [Accessed 07-01-2024]. 2016.
- [3] Erhan undefinedinlar. *Probability and Stochastics*. Springer New York, 2011. ISBN: 9780387878591. DOI: [10.1007/978-0-387-87859-1](https://doi.org/10.1007/978-0-387-87859-1). URL: <http://dx.doi.org/10.1007/978-0-387-87859-1>.
- [4] Wikipedia. *Hoeffding's lemma* — *Wikipedia, The Free Encyclopedia*. <http://en.wikipedia.org/w/index.php?title=Hoeffding's%20lemma&oldid=1114715065>. [Online; accessed 04-January-2024]. 2024.
- [5] Wikipedia. *Rényi entropy* — *Wikipedia, The Free Encyclopedia*. <http://en.wikipedia.org/w/index.php?title=R%C3%A9nyi%20entropy&oldid=1190869396>. [Online; accessed 05-January-2024]. 2024.
- [6] Wikipedia. *Vitali set* — *Wikipedia, The Free Encyclopedia*. <http://en.wikipedia.org/w/index.php?title=Vitali%20set&oldid=1187241923>. [Online; accessed 24-December-2023]. 2023.