Statistical Learning Theory Notes

Nong Minh Hieu[1]

[1] School of Physical and Mathematical Sciences, Nanyang Technological University (NTU - Singapore)

Contents

1 Probability settings

1.1 Classification problem

Definition 1.1 (Classifier (h)). ────────────────────────────────
*In **classification problems**, we consider pairs (x, y) where $x \in \mathcal{X}$ and $y \in \mathcal{Y}$. Where:*

- *$\mathcal{X}$ is the space of **feature vectors**.*

- *$\mathcal{Y}$ is the space of **labels**.*

A classifier is a function $h : \mathcal{X} \to \mathcal{Y}$ which aims to assign correct labels to given feature vectors.

Remark : The key assumptions of classification problems are:

- There exists a joint distribution P_{XY} on $\mathcal{X} \times \mathcal{Y}$.

- The pairs (x, y) (observed data) are random samples of the random variables pair (X, Y) which has the distribution P_{XY}.

Definition 1.2 (Decomposition of P_{XY}). ────────────────────────────
We can decompose P_{XY} in either of the following two ways:

$$P_{XY} = P_{X|Y} P_Y$$
$$P_{XY} = P_{Y|X} P_X$$

Which can be understood as two possible ways to generate the pairs (x, y) from the joint distribution P_{XY}.

- *The first way is to generate a random label $y \sim P_Y$. Then, generate the feature vector corresponding to that label $x \sim P_{X|Y=y}$.*

- *The second way is to generate a random vector $x \sim P_X$. Then, generate the label corresponding to that feature vector $y \sim P_{Y|X=x}$.*

Proposition 1.1: Law of total expectation

Given $\phi : \mathcal{X} \times \mathcal{Y} \to \mathbb{R}$. The **law of total expectation** states that:

$$\mathbb{E}_{XY}\Big[\phi(X, Y)\Big] = \mathbb{E}_Y\Big[\mathbb{E}_{X|Y}[\phi(X, Y)]\Big]$$
$$= \mathbb{E}_X\Big[\mathbb{E}_{Y|X}[\phi(X, Y)]\Big]$$

Similar to how P_{XY} is decomposed, law of total expectation describes two way of taking the average value:

- Loop through the labels and take average over the feature vectors corresponding to each label.

- Loop through the feature vectors and take average over the labels corresponding to each vector.

Proof (Proposition 1.1). ───────────────────────────────
We have:

$$\mathbb{E}_{XY}\Big[\phi(X,Y)\Big] = \int_{\mathcal{X}}\int_{\mathcal{Y}}\phi(x,y)P_{XY}(x,y)dydx$$

$$= \int_{\mathcal{X}}\int_{\mathcal{Y}}\phi(x,y)P_X(x)P_{Y|X}(y|x)dydx$$

$$= \int_{\mathcal{X}}P_X(x)\int_{\mathcal{Y}}\phi(x,y)P_{Y|X}(y|x)dydx$$

$$= \int_{\mathcal{X}}P_X(x)\mathbb{E}_{Y|X=x}\Big[\phi(X,Y)\Big]dx$$

$$= \mathbb{E}_X\Big[\mathbb{E}_{Y|X}\Big[\phi(X,Y)\Big]\Big]$$

Applying the same technique, we have $\mathbb{E}_{XY}\Big[\phi(X,Y)\Big] = \mathbb{E}_Y\Big[\mathbb{E}_{X|Y}[\phi(X,Y)]\Big]$. □.

Remark : Usually, the label space is discrete and finite, meaning $\mathcal{Y} = \{0,1,2,\ldots,m\}$ for some $m < \infty$. Hence, the expectations over Y can be written as discrete sums:

$$\mathbb{E}_{XY}\Big[\phi(X,Y)\Big] = \mathbb{E}_Y\Big[\mathbb{E}_{X|Y}[\phi(X,Y)]\Big] = \sum_{y\in\mathcal{Y}}\mathbb{E}_{X|Y=y}[\phi(X,Y)]$$

$$= \mathbb{E}_X\Big[\mathbb{E}_{Y|X}[\phi(X,Y)]\Big] = \mathbb{E}_X\left[\sum_{y\in\mathcal{Y}}\mathbb{E}_{Y=y|X}[\phi(X,Y)]\right]$$

Definition 1.3 (Hypothesis space ($\mathcal{H}$)). ───────────────────────────
The hypothesis space is a collection (family) of classifiers $h : \mathcal{X} \to \mathcal{Y}$ that have some common properties:

$$\mathcal{H} = \Big\{h : \mathcal{X} \to \mathcal{Y}\Big|\text{some common properties}\Big\}$$

For example, let $\mathcal{X} = \mathbb{R}^d, \mathcal{Y} = (0,1)$. In logistic regression, we assume the classifiers to be logit functions:

$$\mathcal{H}_{logit} = \Big\{h : \mathbb{R}^d \to (0,1)\Big|h(x) = logit(\beta x) = \frac{1}{1 + e^{-\beta x}}, \beta \in \mathbb{R}^{1\times d}\Big\}$$

Definition 1.4 (Learning algorithm ($\mathcal{L}_n$)). ───────────────────────────
*To learn a classifier $h : \mathcal{X} \to \mathcal{Y}$, suppose that we have access to a training dataset of n data pairs $\{(X_k,Y_k)\}_{k=1}^n$ which are assumed to be **i.i.d sampled from** P_{XY}. The domain of the training data is then $(\mathcal{X}\times\mathcal{Y})^n$. A **learning algorithm**, denoted as $\mathcal{L}_n$ is a function/procedure that derives a classifier $\hat{h}_n : \mathcal{X} \to \mathcal{Y}$ from the training data.*

$$\mathcal{L}_n : (\mathcal{X}\times\mathcal{Y})^n \to \mathcal{H}$$

$$\hat{h}_n = \mathcal{L}_n((X_1,Y_1),\ldots,(X_n,Y_n))$$

1.2 Goal of classification

Definition 1.5 (Risk ($R(h)$)). ───────────────────────────
*The **risk** of a classifier is defined as followed:*

$$R(h) = P(h(X) \neq Y) = \mathbb{E}[\mathbf{1}_{\{h(X) \neq Y\}}]$$

Where (X, Y) are independent of the training data.

Definition 1.6 (Bayes Risk (R^*)). ───────────────────────────
*The **Bayes risk** is the infimum of the risk taken over all $h : \mathcal{X} \to \mathcal{Y}$, not just for $h \in \mathcal{H}$:*

$$R^* = \inf_{h:\mathcal{X} \to \mathcal{Y}} R(h)$$

Definition 1.7 (Consistency of learning algorithms). ───────────────────────────
A learning algorithm $\mathcal{L}_n$ is called:

- ***Weakly consistent** if $R(\hat{h}_n) \xrightarrow{p} R^*$:*

  $$\lim_{n \to \infty} P(R(\hat{h}_n) \leq r) = P(R^* \leq r), \ \forall r \geq 0$$

- ***Strongly consistent** if $R(\hat{h}_n) \xrightarrow{a.s} R^*$:*

  $$P\left(\lim_{n \to \infty} \left| R(\hat{h}_n) - R^* \right| \geq \epsilon \right) = 0, \ \forall \epsilon > 0$$

- ***Universally weakly/strongly consistent** if $\mathcal{L}_n$ is weakly/strongly consistent for all P_{XY}. Meaning, consistency holds without any assumption about P_{XY}.*

2 Bayes classifier

2.1 Properties of Bayes Risk

Overview : Recall that the Bayes classifier is the one with minimum risk and the corresponding risk is called the Bayes Risk. For $\mathcal{Y} = \{0,1\}$ and defined:

$$\eta(x) = P(Y = 1 | X = x)$$

Define the following classifier:

$$h^*(x) = \begin{cases} 1 & \text{if } \eta(x) \geq \frac{1}{2} \\ 0 & \text{otherwise} \end{cases}$$

Theorem 2.1: Properties of Bayes classifier

The following properties hold for the Bayes classifier with $\mathcal{Y} = \{0,1\}$ (Binary classification):

- (i) $R(h^*) = \inf_{h:\mathcal{X} \to \mathcal{Y}}\{R(h)\} = R^*$.

- (ii) $\underbrace{R(h) - R^*}_{\text{Excess risk}} = 2\mathbb{E}_X\left[\left|\eta(x) - \frac{1}{2}\right|\mathbf{1}_{\{h(X) \neq h^*(X)\}}\right]$.

- (iii) $R^* = \mathbb{E}\left[\min(\eta(X), 1 - \eta(x))\right]$.

Proof (Theorem 2.1). ———————————————————
Proving each point:

(i) $R(h^*) = \inf_{h:\mathcal{X} \to \mathcal{Y}}\{R(h)\} = R^*$.
For all $h : \mathcal{X} \to \mathcal{Y}$, we have:

$$R(h) = \mathbb{E}_{XY}\left[\mathbf{1}_{\{h(X) \neq Y\}}\right]$$

$$= \mathbb{E}_{x \sim X}\left[\mathbb{E}_{Y|X=x}\left[\mathbf{1}_{\{Y \neq h(x)\}}\right]\right]$$

$$= \mathbb{E}_{x \sim X}\left[\sum_{y \in \{0,1\}} \mathbf{1}_{\{y \neq h(x)\}}\right]$$

$$= \mathbb{E}_{x \sim X}\left[\eta(x)\mathbf{1}_{\{h(x)=0\}} + (1 - \eta(x))\mathbf{1}_{\{h(x)=1\}}\right]$$

Since the two events $\{h(x) = 1\}$ and $\{h(x) = 0\}$ are mutually exclusive, $R(h)$ is the smallest when we set $h(x) = 1$ when $\eta(x) \geq 1 - \eta(x) \implies \eta(x) \geq \frac{1}{2}$. Therefore, we have:

$$h^*(x) = \begin{cases} 1 & \text{if } \eta(x) \geq \frac{1}{2} \\ 0 & \text{otherwise} \end{cases}$$

(ii) $\underbrace{R(h) - R^*}_{\textit{Excess risk}} = 2\mathbb{E}_X\left[\left|\eta(x) - \frac{1}{2}\right|\mathbf{1}_{\{h(X) \neq h^*(X)\}}\right]$.

We have:

$$R(h) - R^* = \mathbb{E}_{x \sim X}\left[\mathbb{E}_{Y|X=x}\left[\mathbf{1}_{\{Y \neq h(x)\}}\right]\right] - \mathbb{E}_{x \sim X}\left[\mathbb{E}_{Y|X=x}\left[\mathbf{1}_{\{Y \neq h^*(x)\}}\right]\right]$$

$$= \mathbb{E}_{x \sim X}\left[\sum_{y \in \{0,1\}} \mathbf{1}_{\{y \neq h(x)\}} P(Y = y | X = x)\right] - \mathbb{E}_{x \sim X}\left[\sum_{y \in \{0,1\}} \mathbf{1}_{\{y \neq h^*(x)\}} P(Y = y | X = x)\right]$$

$$= \mathbb{E}_{x \sim X}\left[\eta(x)\left(\mathbf{1}_{\{h(x)=0\}} - \mathbf{1}_{\{h^*(x)=0\}}\right) + (1 - \eta(x))\left(\mathbf{1}_{\{h(x)=1\}} - \mathbf{1}_{\{h^*(x)=1\}}\right)\right]$$

$$= \mathbb{E}_{x \sim X}\left[\eta(x)\left(\mathbf{1}_{\{h(x) \neq h^*(x), h(x)=0\}} - \mathbf{1}_{\{h(x) \neq h^*(x), h(x)=1\}}\right)\right.$$

$$\left. + (1 - \eta(x))\left(\mathbf{1}_{\{h(x) \neq h^*(x), h(x)=1\}} - \mathbf{1}_{\{h(x) \neq h^*(x), h(x)=0\}}\right)\right]$$

$$= \mathbb{E}_{x \sim X}\left[(2\eta(x) - 1)\mathbf{1}_{\{h(x) \neq h^*(x), h(x)=0\}} + (1 - 2\eta(x))\mathbf{1}_{\{h(x) \neq h^*(x), h(x)=1\}}\right]$$

$$= \mathbb{E}_{x \sim X}\left[\left|2\eta(x) - 1\right|\mathbf{1}_{\{h(x) \neq h^*(x)\}}\right]$$

$$= 2\mathbb{E}_X\left[\left|\eta(X) - \frac{1}{2}\right|\mathbf{1}_{\{h(X) \neq h^*(X)\}}\right]$$

(iii) $R^* = \mathbb{E}\left[\min(\eta(X), 1 - \eta(x))\right].$
From (i) we have:

$$R(h^*) = \mathbb{E}_{x \sim X}\left[\eta(x)\mathbf{1}_{\{h^*(x)=0\}} + (1 - \eta(x))\mathbf{1}_{\{h^*(x)=1\}}\right]$$

$$= \mathbb{E}_X\left[\min(\eta(X), 1 - \eta(x))\right]$$

$\square$.

> **Theorem 2.2: Properties of Bayes classifier (Multi-class)**
>
> For multi-class classification with more than two labels : $\mathcal{Y} = \{1, 2, \ldots, M\}$, the Bayes classifier is defined as followed:
>
> $$h^*(x) = \arg\max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\}$$
>
> $$\text{Where} : \eta_y(x) = P(Y = y | X = x)$$
>
> The following properties hold for the Bayes classifier with $\mathcal{Y} = \{1, 2, \ldots, M\}$ (Multi-class classification):
>
> - (i) **Bayes Risk** R^* :
>
> $$R^* = \mathbb{E}_{x \sim X} \left[1 - \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} \right] = \mathbb{E}_{x \sim X} \left[\min_{y \in \mathcal{Y}} \overline{\eta_y}(x) \right]$$
>
> - (ii) **Excess Risk** $R(h) - R^*$:
>
> $$R(h) - R^* = \mathbb{E}_X \left[\left(\eta_{y_x^*}(x) - \eta_{y_x}(x) \right) \mathbf{1}_{\{h(x) \neq h^*(x)\}} \right]$$
>
> Where $y_x = h(x)$ is the prediction made by an arbitrary classifier $h : \mathcal{X} \to \mathcal{Y}$ and $y_x^* = h^*(x)$ is the prediction made by the Bayes classifier.

Proof (Theorem 2.2). ───────────────────────────────────
(The proof of this theorem has been included in the solution of Exercise 2.1). □.

2.2 Likelihood Ratio Test

Overview : Define $\pi_1 = P(Y = 1)$ and $\pi_0 = P(Y = 0)$ be the prior probabilities. Let $p_1(x) = P(X = x | Y = 1)$ and $p_0(x) = P(X = x | Y = 0)$ be the class-conditional densities. Note that we have:

$$
\begin{aligned}
\eta(x) &= P(Y = 1 | X = x) \\
&= \frac{P(X = x | Y = 1)P(Y = 1)}{P(X = x | Y = 1)P(Y = 1) + P(X = x | Y = 0)P(Y = 0)} \\
&= \frac{\pi_1 p_1(x)}{\pi_1 p_1(x) + \pi_0 p_0(x)} \\
&= \frac{1}{1 + \frac{\pi_0 p_0(x)}{\pi_1 p_1(x)}}
\end{aligned}
$$

Hence, we have:

$$
\begin{aligned}
\eta(x) \geq \frac{1}{2} &\iff \frac{\pi_0 p_0(x)}{\pi_1 p_1(x)} \\
&\iff \frac{p_1(x)}{p_0(x)} \geq \frac{\pi_0}{\pi_1}
\end{aligned}
$$

> **Proposition 2.1: Likelihood ratio test**
>
> The Bayes classifier h^* can be re-defined as followed:
>
> $$h^*(x) = \begin{cases} 1 & \text{if } \frac{p_1(x)}{p_0(x)} \geq \frac{\pi_0}{\pi_1} \\ \\ 0 & \text{otherwise} \end{cases}$$
>
> The fraction $\frac{p_1(x)}{p_0(x)}$ is called the **likelihood ratio**.

2.3 Plug-in classifier

Definition 2.1 (Plug-in classifier). ───────────────
*A **plug-in classifier** is based on an estimate of $\eta(x)$. This estimate is then plugged into the definition of the Bayes classifier. Suppose that $\widehat{\eta_n}$ is an estimate of η based on n training samples $\{(X_i, Y_i)\}_{i=1}^n$. We define $\widehat{h_n}$ as:*

$$\widehat{h_n} = \begin{cases} 1 & \text{if } \widehat{\eta_n}(x) \geq \frac{1}{2} \\ \\ 0 & \text{otherwise} \end{cases}$$

> **Corollary 2.1: Excess risk of plug-in classifier**
>
> We have the following upper-bound for the excess risk of the plug-in classifier:
>
> $$R(\widehat{h_n}) - R^* \leq 2\mathbb{E}_X\left[\left|\eta(X) - \widehat{\eta_n}(X)\right|\right]$$

Proof (Corollary 2.1). ───────────────
From theorem 2.1, we have:

$$R(\widehat{h_n}) - R^* = 2\mathbb{E}_X\left[\left|\eta(X) - \frac{1}{2}\right|\mathbf{1}_{\{\widehat{h_n}(X) \neq h^*(X)\}}\right]$$

The indicator term will be non-zero in the above equality if one of the following cases occurs:

$$\begin{cases} \widehat{h_n}(X) = 1, h^*(X) = 0 \\ \\ \widehat{h_n}(X) = 0, h^*(X) = 1 \end{cases} \implies \begin{cases} \widehat{\eta_n}(X) \geq \frac{1}{2}, \eta(X) < \frac{1}{2} \\ \\ \widehat{\eta_n}(X) < \frac{1}{2}, \eta(X) \geq \frac{1}{2} \end{cases}$$

Case 1 : $\widehat{\eta_n}(X) \geq \frac{1}{2}, \eta(X) < \frac{1}{2}$
We have:

$$\eta(X) - \widehat{\eta_n}(X) \leq \eta(X) - \frac{1}{2} \quad (\text{Both sides negative})$$

$$\implies \left|\eta(X) - \widehat{\eta_n}(X)\right| \geq \left|\eta(X) - \frac{1}{2}\right|$$

Case 2 : $\widehat{\eta_n}(X) < \frac{1}{2}, \eta(X) \geq \frac{1}{2}$

We have:

$$\widehat{\eta_n}(X) - \eta(X) \geq \widehat{\eta_n}(X) - \frac{1}{2} \geq \eta(X) - \frac{1}{2} \quad \text{(All positive)}$$

Therefore, we have:

$$\left| \eta(X) - \widehat{\eta_n}(X) \right| \geq \left| \eta(X) - \frac{1}{2} \right|$$

For both cases, we have the same $\left| \eta(X) - \widehat{\eta_n}(X) \right| \geq \left| \eta(X) - \frac{1}{2} \right|$ *inequality. Therefore, we have:*

$$R(\widehat{h_n}) - R^* \leq 2\mathbb{E}_X \left[\left| \eta(X) - \widehat{\eta_n}(X) \right| \right]$$

$\square$.

2.4   End of chapter exercises

Exercise 2.1

Extend theorem 2.1 to the multi-class classification case where $\mathcal{Y} = \{1, 2, \ldots, M\}$. In other words, prove theorem 2.2.

Solution (Exercise 2.1).

We re-define the Bayes classifier h^* as followed:

$$h^*(x) = \arg\max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\},$$

$$\eta_y(x) = P(Y = y | X = x)$$

We have:

$$\sum_{y \in \mathcal{Y}} \eta_y(x) = 1, \ \forall x \in \mathcal{X}$$

(i) *Calculate Bayes risk* R^*

For any classifier $h : \mathcal{X} \to \mathcal{Y}$, we have:

$$R(h) = \mathbb{E}_{x \sim X} \left[\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) \right]$$

Letting $\hat{y}_x = h(x)$ being h's prediction for a given feature vector $x \in \mathcal{X}$, we have:

$$R(h) = \mathbb{E}_{x \sim X} \left[\sum_{y \in \mathcal{Y}; y \neq \hat{y}_x} \eta_y(x) \right] = \mathbb{E}_{x \sim X} \left[1 - \eta_{\hat{y}_x}(x) \right]$$

In order to minimize $R(h)$, we need $\eta_{\hat{y}_x}(x)$ to be maxmized for all $x \in \mathcal{X}$. Hence, we have:

$$R^* = \mathbb{E}_{x \sim X} \left[1 - \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} \right]$$

Therefore, we have $h^*(x) = \arg\max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\}$ is the Bayes classifier and the Bayes risk $R^* = \mathbb{E}_{x \sim X} \left[1 - \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} \right]$.

(ii) *Calculate excess risk* $R(h) - R^*$

For any $h : \mathcal{X} \to \mathcal{Y}$, we have:

$$R(h) - R^* = \mathbb{E}_{x \sim X} \left[\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) \right] - \mathbb{E}_{x \sim X} \left[1 - \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} \right]$$

$$= \mathbb{E}_{x \sim X} \left[\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) + \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} - 1 \right]$$

Denote $h^*(x) = y_x^*$ and $h(x) = y_x$. When $h(x) = h^*(x) = y_x^*$, we have:

$$\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) + \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} = \sum_{y \in \mathcal{Y}; y \neq y_x} \eta_y(x) + \eta_{y_x^*}(x)$$

$$= \sum_{y \in \mathcal{Y}; y \neq y_x^*} \eta_y(x) + \eta_{y_x^*}(x)$$

$$= \sum_{y \in \mathcal{Y}} \eta_y(x) = 1$$

$$\implies \sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) + \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} - 1 = 0$$

When $h(x) \neq h^*(x)$, we have:

$$\sum_{y \in \mathcal{Y}} \mathbf{1}_{\{h(x) \neq y\}} \eta_y(x) + \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} - 1 = \sum_{y \in \mathcal{Y}; y \neq y_x} \eta_y(x) + \eta_{y_x^*}(x) - 1$$

$$= 2\eta_{y_x^*}(x) - 1 + \sum_{y \in \mathcal{Y} \setminus \{y_x, y_x^*\}} \eta_y(x)$$

$$= 2\eta_{y_x^*}(x) - \left(\eta_{y_x}(x) + \eta_{y_x^*}(x) \right)$$

$$= \eta_{y_x^*}(x) - \eta_{y_x}(x).$$

Therefore, we can re-write the excess risk by multiplying the entire integrand with the indicator function $\mathbf{1}_{\{h(x) \neq h^*(x)\}}$ as followed:

$$R(h) - R^* = \mathbb{E}_{x \sim X} \left[\left(\eta_{y_x^*}(x) - \eta_{y_x}(x) \right) \mathbf{1}_{\{h(x) \neq h^*(x)\}} \right]$$

(iii) *Simpler form of Bayes risk*
From (i) we have:

$$R^* = \mathbb{E}_X \left[1 - \max_{y \in \mathcal{Y}} \left\{ \eta_y(x) \right\} \right] = \mathbb{E}_X \left[\min_{y \in \mathcal{Y}} \left\{ \overline{\eta_y}(x) \right\} \right]$$

Where $\overline{\eta_y}(x) = P(Y \neq y | X = x)$.

$\square$.

Exercise 2.2

Define the α-**cost-sensitive risk** of a classifier $h : \mathcal{X} \to \mathcal{Y}$ as followed:

$$R_\alpha(h) = \mathbb{E}_{XY} \left[(1 - \alpha) \mathbf{1}_{\{Y=1, h(X)=0\}} + \alpha \mathbf{1}_{\{Y=0, h(X)=1\}} \right]$$

Define the Bayes classifier and prove and analogue of theorem 2.1.

Solution (Exercise 2.2).
Using the law of total expectation, we have:

$$R_\alpha(h) = \mathbb{E}_{x \sim X} \left[\sum_{y \in \{0,1\}} \left[(1 - \alpha) \mathbf{1}_{\{y=1, h(x)=0\}} + \alpha \mathbf{1}_{\{y=0, h(x)=1\}} \right] P(Y = y | X = x) \right]$$

$$= \mathbb{E}_{x \sim X} \left[(1 - \alpha) \eta(x) \mathbf{1}_{\{h(x)=0\}} + \alpha (1 - \eta(x)) \mathbf{1}_{\{h(x)=1\}} \right]$$

Since $\mathbf{1}_{\{h(x)=0\}}$ and $\mathbf{1}_{\{h(x)=1\}}$ are mutually exclusive, in order for $R_\alpha(h)$ to be minimize, we define the following Bayes classifier:

$$h^*(x) = \begin{cases} 1 & \text{if } \alpha(1-\eta(x)) \leq (1-\alpha)\eta(x) \\ 0 & \text{otherwise} \end{cases} = \begin{cases} 1 & \text{if } \eta(x) \geq \alpha \\ 0 & \text{otherwise} \end{cases}$$

We can also derive a likelihood-ratio test version of the Bayes classifier, we have:

$$\eta(x) \geq \alpha \implies \frac{1}{1 + \frac{\pi_0 p_0(x)}{\pi_1 p_1(x)}} \geq \alpha$$

$$\implies 1 + \frac{\pi_0 \cdot p_0(x)}{\pi_1 \cdot p_1(x)} \leq \frac{1}{\alpha}$$

$$\implies \frac{p_1(x)}{p_0(x)} \geq \frac{\alpha}{1-\alpha} \cdot \frac{\pi_0}{\pi_1}$$

Hence, we can rewrite the Bayes classifier as followed:

$$h^*(x) = \begin{cases} 1 & \text{if } \frac{p_1(x)}{p_0(x)} \geq \frac{\alpha}{1-\alpha} \cdot \frac{\pi_0}{\pi_1} \\ 0 & \text{otherwise} \end{cases}$$

(i) **Bayes Risk** R_α^*
We have:

$$\begin{aligned} R_\alpha^* &= R_\alpha(h^*) \\ &= \mathbb{E}_{x \sim X}\Big[(1-\alpha)\eta(x)\mathbf{1}_{\{h^*(x)=0\}} + \alpha(1-\eta(x))\mathbf{1}_{\{h^*(x)=1\}}\Big] \\ &= \mathbb{E}_X\Big[\min(\alpha(1-\eta(X)), (1-\alpha)\eta(X))\Big] \end{aligned}$$

(ii) **Excess Risk** $R_\alpha(h) - R_\alpha^*$
For an arbitrary $h : \mathcal{X} \to \mathcal{Y}$, we have:

$$\begin{aligned} R_\alpha(h) - R_\alpha^* &= \mathbb{E}_{x \sim X}\Big[(1-\alpha)\eta(x)\Big(\mathbf{1}_{\{h(x)=0\}} - \mathbf{1}_{\{h^*(x)=0\}}\Big) + \alpha(1-\eta(x))\Big(\mathbf{1}_{\{h(x)=1\}} - \mathbf{1}_{\{h^*(x)=1\}}\Big)\Big] \\ &= \mathbb{E}_{x \sim X}\Big[(1-\alpha)\eta(x)\Big(\mathbf{1}_{\{h(x)=0,h^*(x)=1\}} - \mathbf{1}_{\{h(x)=1,h^*(x)=0\}}\Big) \\ &\qquad + \alpha(1-\eta(x))\Big(\mathbf{1}_{\{h(x)=1,h^*(x)=0\}} - \mathbf{1}_{\{h(x)=0,h^*(x)=1\}}\Big)\Big] \\ &= \mathbb{E}_{x \sim X}\Big[\mathbf{1}_{\{h(x)=0,h^*(x)=1\}}(\eta(x) - \alpha) + \mathbf{1}_{\{h(x)=1,h^*(x)=0\}}(\alpha - \eta(x))\Big] \\ &= \mathbb{E}_X\Big[\big|\eta(X) - \alpha\big|\mathbf{1}_{\{h(X)\neq h^*(X)\}}\Big] \end{aligned}$$

$\square$.

3 Hoeffding's inequality

3.1 Markov's Inequality

> **Proposition 3.1: Markov's Inequality**
>
> Let U be a non-negative random variable on $\mathbb{R}$, then for all $t > 0$, we have:
> $$P(U \geq t) \leq \frac{1}{t}\mathbb{E}[U]$$

Proof (Proposition 3.1).
We have:

$$
\begin{aligned}
tP(U \geq t) &= t\mathbb{E}\Big[\mathbf{1}_{\{U \geq t\}}\Big] \\
&= t\int_0^\infty \mathbf{1}_{\{x \geq t\}} f_U(x)dx \\
&= t\int_t^\infty f_U(x)dx \\
&\leq \int_t^\infty x f_U(x)dx \\
&\leq \int_0^\infty x f_U(x)dx = \mathbb{E}[U] \\
\implies P(U \geq t) &\leq \frac{1}{t}\mathbb{E}[U]
\end{aligned}
$$

$\square$.

> **Corollary 3.1: Chebyshev's Inequality**
>
> Let Z be a random variable on $\mathbb{R}$ with mean μ and variance σ^2, we have:
> $$P\Big(\big|Z - \mu\big| \geq t\Big) \leq \frac{\sigma^2}{t^2}$$

Proof (Corollary 3.1).
Using Markov's inequality, we have:

$$
\begin{aligned}
P\Big(\big|Z - \mu\big| \geq t\Big) &= P\Big(\big|Z - \mu\big|^2 \geq t^2\Big) \\
&\leq \frac{\mathbb{E}\Big[\big|Z - \mu\big|^2\Big]}{t^2} = \frac{\sigma^2}{t^2}
\end{aligned}
$$

$\square$.

> **Corollary 3.2: Chernoff's bounding method**
>
> Let Z be a random variable on $\mathbb{E}$, for any $t > 0$, we have:
> $$P(Z \geq t) \leq \inf_{s>0} e^{-st} M_Z(s)$$

Proof (Corollary 3.2).

We have:

$$P(Z \geq t) = P(sZ \geq st), \quad (t > 0)$$
$$= P(e^{sZ} \geq e^{st})$$
$$\leq \frac{\mathbb{E}\left[e^{sZ}\right]}{e^{st}} = e^{-st} M_Z(s) \quad (\text{Markov's inequality})$$

Since the above inequality holds for all $s > 0$, we can just take the infimum to obtain the tightest bound. Hence, we have:

$$P(Z \geq t) \leq \inf_{s>0} e^{-st} M_Z(s)$$

$\square$.

3.2 Hoeffding's Inequality

Before diving into Hoeffding's inequality, we need to go through the following lemma (whose proof will not be included) that will help us prove the Hoeffding's inequality:

Lemma 3.1: Hoeffding's lemma

Let V be a random variable on $\mathbb{R}$ with $\mathbb{E}[V] = 0$ and suppose that $a \leq V \leq b$ with probability one. We have:

$$\mathbb{E}\left[e^{sV}\right] \leq \exp\left(\frac{s^2(b-a)^2}{8}\right)$$

Proof (Lemma 3.1). ——————————————————
(The proof for this lemma can be found here [3]). $\square$.

Theorem 3.1: Hoeffding's Inequality

Let $Z_1, Z_2, \ldots, Z_n$ be independent random variables on $\mathbb{R}$ such that $a_i \leq Z_i \leq b_i$ with probability one for all $1 \leq i \leq n$. Let $S_n = \sum_{i=1}^{n} Z_i$. We have:

$$P\left(\left|S_n - \mathbb{E}[S_n]\right| \geq t\right) \leq 2\exp\left(-\frac{2t^2}{\sum_{i=1}^{n}(b_i - a_i)^2}\right), \quad \forall t > 0$$

Proof (Theorem 3.1). ——————————————————

Using the Chernoff's bounds, we have:

$$P\Big(\Big|S_n - \mathbb{E}[S_n]\Big| \geq t\Big) \leq \inf_{s>0} e^{-st} M_{S_n - \mathbb{E}[S_n]}(s)$$

$$= \inf_{s>0} e^{-st} \mathbb{E}\Big[e^{s(S_n - \mathbb{E}[S_n])}\Big]$$

$$= \inf_{s>0} e^{-st} \mathbb{E}\Big[\exp\Big(s\sum_{i=1}^n (Z_i - \mathbb{E}[Z_i])\Big)\Big]$$

$$= \inf_{s>0} e^{-st} \mathbb{E}\Big[\prod_{i=1}^n \exp\Big(s(Z_i - \mathbb{E}[Z_i])\Big)\Big]$$

$$= \inf_{s>0} e^{-st} \prod_{i=1}^n \mathbb{E}\Big[\exp\Big(s(Z_i - \mathbb{E}[Z_i])\Big)\Big] \quad (\textit{Since all } Z_i - \mathbb{E}[Z_i] \textit{ are independent})$$

$$\leq \inf_{s>0} e^{-st} \prod_{i=1}^n \exp\Big(\frac{s^2(b_i - a_i)^2}{8}\Big) \quad (\textit{By Hoeffding's lemma})$$

$$= \inf_{s>0} \exp\Big(-st + \sum_{i=1}^n \frac{s^2(b_i - a_i)^2}{8}\Big)$$

In order for the above to be minimized, we differentiate the term inside the exponential and set the derivative to 0 to find the optimal $s > 0$. We have:

$$-t + s\sum_{i=1}^n \frac{(b_i - a_i)^2}{4} = 0 \implies s = \frac{4t}{\sum_{i=1}^n (b_i - a_i)^2}$$

Letting $c = \sum_{i=1}^n (b_i - a_i)^2$, we now can derive the tightest Chernoff's bound as followed:

$$P\Big(\Big|S_n - \mathbb{E}[S_n]\Big| \geq t\Big) \leq \exp\Big(-\frac{4t^2}{c} + \frac{16t^2}{c^2}\cdot\frac{c}{8}\Big) = \exp\Big(-\frac{2t^2}{c}\Big)$$

$$= \exp\Big(-\frac{2t^2}{\sum_{i=1}^n (b_i - a_i)^2}\Big)$$

$\square$.

3.3   Convergence of Empirical Risk

Definition 3.1 (Empirical Risk). ─────────────────────────────────

*Suppose we are given training data $\Big\{(X_i, Y_i)_{i=1}^n\Big\}$ such that each pair $(X_i, Y_i) \sim P_{XY}$ are independently identically distributed. Let $h : \mathcal{X} \to \mathcal{Y}$ be a classifier. We define the **empirical risk** to be:*

$$\widehat{R_n}(h) = \frac{1}{n}\sum_{i=1}^n \mathbf{1}_{\{h(X_i) \neq Y_i\}}$$

*Note that $\mathbb{E}[\widehat{R_n}(h)] = R(h)$ and $n\widehat{R_n}(h) \sim Binomial(n, R(h))$. In the following corollary of the Hoeffding's inequality, we will answer the question **how close the empirical risk is as an estimate of true risk** or **how fast the empirical risk converges to the true risk**.*

> **Corollary 3.3: Convergence of Empirical Risk**
>
> Given training data $\left\{(X_i, Y_i)_{i=1}^n\right\}$ such that each pair $(X_i, Y_i) \sim P_{XY}$ are independently identically distributed. Let $h : \mathcal{X} \to \mathcal{Y}$ be a classifier, we have:
>
> $$P\left(\left|\widehat{R_n}(h) - R(h)\right| \geq \epsilon\right) \leq 2e^{-2n\epsilon^2}, \quad \epsilon > 0$$

Proof (Corollary 3.3). _____

For all $1 \leq i \leq n$, we have $\mathbf{1}_{\{h(X_i) \neq Y_i\}} \in \{0, 1\}$. Hence, with probability one, $0 \leq \mathbf{1}_{\{h(X_i) \neq Y_i\}} \leq 1$ and $b_i = 1, a_i = 0$ for all $1 \leq i \leq n$.

Using the Hoeffding's inequality, we have:

$$
\begin{aligned}
P\left(\left|\widehat{R_n}(h) - R(h)\right| \geq \epsilon\right) &= P\left(\left|\widehat{R_n}(h) - \mathbb{E}[\widehat{R_n}(h)]\right| \geq \epsilon\right) \\
&= P\left(\left|n\widehat{R_n}(h) - \mathbb{E}[n\widehat{R_n}(h)]\right| \geq n\epsilon\right) \\
&\leq \exp\left(-\frac{2n^2\epsilon^2}{\sum_{i=1}^n (b_i - a_i)^2}\right) \quad \text{(Hoeffding's inequality)} \\
&= e^{-2n\epsilon^2}
\end{aligned}
$$

$\square$.

3.4   KL-divergence & Hypothesis Testing

Set-up (Hypothesis Testing) : Suppose that we have $\mathcal{Y} = \{0, 1\}$ and P_{XY} is a distribution on $\mathcal{X} \times \mathcal{Y}$. Let's assume that:

- The prior probabilities π_y are equal.

- The supports of likelihoods p_0, p_1 are the same.

- $0 < \alpha \leq p_y(x) \leq \beta < \infty$ for all $x \in \mathcal{X}$ such that $p_y(x) > 0$ and for all $y \in \{0, 1\}$.

Now suppose $X_1, \ldots, X_n \sim p_y$ are independently identically distributed where $y \in \{0, 1\}$ is unknown. Can we guess y and how good our guess would be?

Proposition 3.2: KL-divergence hypothesis testing

From the above settings, the optimal classifier is given by the likelihood ratio test:

$$\widehat{h_n}(x) = \begin{cases} 1 & \text{if } \frac{\prod_{i=1}^n p_1(x_i)}{\prod_{i=1}^n p_0(x_i)} \geq \frac{\pi_0}{\pi_1} \quad (=1) \\ \\ 0 & \text{otherwise} \end{cases}$$

Where $x = \left(x_1, \ldots, x_n \right)$ is an observation of the random vector $X = \left(X_1, \ldots, X_n \right)$. Define the class-specific risk $R_y(h)$ be the risk of misclassification when the true label is $Y = y$:

$$R_y(h) = P(h(X) \neq Y | Y = y)$$

Then, we have:

$$R_0(\widehat{h_n}) \leq e^{-2nD(p_0||p_1)^2/c}, \text{ where } c = 4(\log \beta - \log \alpha)^2$$

Where $D(p_0||p_1)$ is the KL-divergence of p_1 from p_0. We can prove a similar exponentially decaying bound for $R_1(\widehat{h_n})$.

Proof. _____

Proposition 3.2 We can rewrite the optimal classifier as:

$$\widehat{h_n}(X) = \begin{cases} 1 & \text{if } \widehat{S_n}(X_1, \ldots, X_n) \geq 0 \\ \\ 0 & \text{otherwise} \end{cases}$$

Where we have:

$$\begin{aligned} \widehat{S_n}(X_1, \ldots, X_n) &= \log \frac{\prod_{i=1}^n p_1(X_i)}{\prod_{i=1}^n p_0(X_i)} \\ &= \sum_{i=1}^n \log \frac{p_1(X_i)}{p_0(X_i)} \\ &= \sum_{i=1}^n Z_i \quad \left(\text{Letting } Z_i = \log \frac{p_1(X_i)}{p_0(X_i)} \right) \end{aligned}$$

Since the likelihoods are bounded, we have:

$$a_i = \log \frac{\alpha}{\beta} \leq Z_i \leq \log \frac{\beta}{\alpha} = b_i, \quad 1 \leq i \leq n$$

Now, we have:

$$\begin{aligned} R_0(\widehat{h_n}) &= P(h(X) \neq Y | Y = 0) \\ &= P(\widehat{S_n} \geq 0 | Y = 0) \\ &= P(\widehat{S_n} - \mathbb{E}[S_n | Y = 0] \geq -\mathbb{E}[S_n | Y = 0] | Y = 0) \end{aligned}$$

To calculate the conditional expectation $\mathbb{E}[S_n | Y = 0]$, we have:

$$\begin{aligned} \mathbb{E}[S_n | Y = 0] &= n\mathbb{E}[Z_1 | Y = 0] \\ &= n \int \log \frac{p_1(x)}{p_0(x)} p_0(x) dx \\ &= -n \int \log \frac{p_0(x)}{p_1(x)} p_0(x) dx = -nD(p_0||p_1) \end{aligned}$$

Therefore, we have:

$$R_0(\widehat{h_n}) = P(\widehat{S_n} - \mathbb{E}[S_n|Y=0] \geq nD(p_0||p_1)|Y=0)$$

$$\leq \exp\left(-\frac{2n^2 D(p_0||p_1)^2}{\sum_{i=1}^n (b_i - a_i)^2}\right) \quad (\textit{Hoeffding's inequality})$$

For every $1 \leq i \leq n$, we have:

$$b_i - a_i = \log\frac{\beta}{\alpha} - \log\frac{\alpha}{\beta}$$

$$= \log\frac{\beta^2}{\alpha^2} = 2\log\frac{\beta}{\alpha} = 2(\log\beta - \log\alpha)$$

$$\implies \sum_{i=1}^n (b_i - a_i)^2 = 4n(\log\beta - \log\alpha)^2$$

Finally, we have:

$$R_0(\widehat{h_n}) \leq \exp\left(-\frac{2nD(p_0||p_1)^2}{4(\log\beta - \log\alpha)^2}\right)$$

Similarly, for $R_1(\widehat{h_n})$, we have:

$$R_1(\widehat{h_n}) \leq \exp\left(-\frac{2nD(p_1||p_0)^2}{4(\log\beta - \log\alpha)^2}\right)$$

$\square$.

3.5 End of chapter exercises

> **Exercise 3.1**
>
> - (**i**) Apply Chernoff's bounding method to obtain an exponential bound on the tail probability $P(Z \geq t)$ for a Gaussian random variable $Z \sim \mathcal{N}(\mu, \sigma^2)$.
>
> - (**ii**) Appealing to the central limit theorem, use part (**i**) to give an approximate bound on the binomial tail. This should not only match the exponential decay given by Hoeffding's inequality, but also reveal the dependence on the variance of the binomial.

Solution (Exercise 3.1).

.

(**i**) *Chernoff's bounds for* $Z \sim \mathcal{N}(\mu, \sigma^2)$
Using the Chernoff's bounding method, we have:

$$P(Z \geq t) \leq \inf_{s>0} e^{-st} M_Z(s)$$

$$= \inf_{s>0} \exp\left(-st + \mu s + \frac{1}{2}\sigma^2 s^2\right)$$

The above bound is the tightest when the derivative of the term inside the exponential equals zero. Hence, we have:

$$-t + \mu + s\sigma^2 = 0 \implies s = \frac{t - \mu}{\sigma^2}$$

From the above, we have the tightest Chernoff's bound as followed:

$$P(Z \geq t) \leq \exp\left(-\frac{(t-\mu)^2}{\sigma^2} + \frac{(t-\mu)^2}{2\sigma^2}\right) = \exp\left(-\frac{(t-\mu)^2}{2\sigma^2}\right)$$

(**ii**) *Binomial tail upper bound*
Let S_n be the binomial random variable such that:

$$S_n = \sum_{i=1}^{n} X_i, \quad X_i \sim Bernoulli(p)$$

For a positive $\epsilon > 0$, we want to know the upper tail bound $P(S_n - \mathbb{E}[S_n] \geq \epsilon)$. Letting $\overline{X} = \frac{1}{n}S_n$, we have:

$$P(S_n - \mathbb{E}[S_n] \geq \epsilon) = P\left(\overline{X} - \frac{\mathbb{E}[S_n]}{n} \geq \frac{\epsilon}{n}\right)$$

$$= P\left(\overline{X} - p \geq \frac{\epsilon}{n}\right)$$

$$= P\left(\frac{\overline{X} - p}{\sqrt{pq}/\sqrt{n}} \geq \frac{\epsilon}{\sqrt{npq}}\right), \quad (q = 1 - p)$$

By the Central Limit Theorem, we have:

$$\frac{\overline{X} - p}{\sqrt{pq}/\sqrt{n}} \xrightarrow{d} \mathcal{N}(0, 1)$$

Hence, as $n \to \infty$, the upper tail bound would be:

$$P(S_n - \mathbb{E}[S_n] \geq \epsilon) = P\left(\frac{\overline{X} - p}{\sqrt{pq}/\sqrt{n}} \geq \frac{\epsilon}{\sqrt{npq}}\right)$$

$$\leq \exp\left(-\frac{\epsilon^2}{2npq}\right) = \exp\left(-\frac{\epsilon^2}{2Var(S_n)}\right)$$

Double-check the bound with Hoeffding's inequality, we have:

$$P(S_n - \mathbb{E}[S_n] \geq \epsilon) \leq \exp\left(-\frac{2\epsilon^2}{n}\right)$$

$\square$.

Exercise 3.2

Can you remove the assumption in $0 < \alpha \leq p_y(x)$? Consider other restrictions on p_y, other concentration inequalities, or other f-divergences.

4 Empirical Risk Minimization

A   List of Definitions

B   Important Theorems

C   Important Corollaries

D   Important Propositions

E References

References

[1] Rick Durrett. *Probability: Theory and Examples*. 4th. USA: Cambridge University Press, 2010. ISBN: 0521765390.

[2] Erhan undefinedinlar. *Probability and Stochastics*. Springer New York, 2011. ISBN: 9780387878591. DOI: 10.1007/978-0-387-87859-1. URL: http://dx.doi.org/10.1007/978-0-387-87859-1.

[3] Wikipedia. *Hoeffding's lemma — Wikipedia, The Free Encyclopedia*. http://en.wikipedia.org/w/index.php?title=Hoeffding's%20lemma&oldid=1114715065. [Online; accessed 04-January-2024]. 2024.

[4] Wikipedia. *Vitali set — Wikipedia, The Free Encyclopedia*. http://en.wikipedia.org/w/index.php?title=Vitali%20set&oldid=1187241923. [Online; accessed 24-December-2023]. 2023.